

Introduction to Database Systems

Lecture 2: Data Models & SQL (Ch. 2.1-2.3)

Data Models

- language / notation for talking about data
- models we will use:
 - relational: data is a collection of tables
 - semi-structured: data is a tree
- other models:
 - key-value pairs: used by NoSQL systems
 - graph data model: used by RDF (semi-structured can also do)
 - object oriented: often layered on relational, J2EE

Relational Model

columns /
attributes /
fields

- Data is a collection of relations / tables:

Name	Country	Employees	For_Profit
GizmoWorks	USA	20000	True
Canon	Japan	50000	True
Hitachi	Japan	30000	True
HappyCam	Canada	500	False

- mathematically, relation is a set of tuples
 - each tuple appears 0 or 1 times in the table
 - order of the rows is unspecified

Relational Schema

- Each column has a “domain” (or type)
 - SQL has Java-like types for numbers, strings, etc.
 - domain is a constraint on the data allowed in the table
- Names and types part of the “schema” of the table:

```
Company(Name: string, Country: string,  
        Employees: int, For_Profit: boolean)
```

- Particular data is an “instance” of that relation
 - data changes over time
 - DBMS usually just stores the current instance

Keys

- Key = subset of columns that uniquely identifies tuple
- Another constraint on the table
 - no two tuples can have the same values for those columns
- Examples:
 - Movie(title, year, length, genre): key is (title, year)
 - what is a good key for Company?
- Part of the schema (book notation is underline):

```
Company(Name: string, Country: string,  
        Employees: int, For_Profit: boolean)
```

Keys (cont.)

- Can have multiple keys for a table
- Only one of those keys may be “primary”
 - DBMS often makes searches by primary key fastest
 - other keys are called “secondary”
- “Foreign key” is a column (or columns) whose value is a key of another table
 - i.e., a reference to another row in another table

SQL (“sequel”)

- Standard query language for relational data
 - used for databases in many different contexts
 - inspires query languages for non-relational (e.g. SQL++)
- Everything not in quotes (‘...’) is case insensitive
- Provides standard types. Examples:
 - numbers: INT, FLOAT, DECIMAL(p,s)
 - DECIMAL(p,s): Exact numerical, precision p, scale s. Example: decimal(5,2) is a number that has 3 digits before the decimal and 2 digits after the decimal
 - strings: CHAR(n), VARCHAR(n)
 - CHAR(n): Fixed-length n
 - VARCHAR(n): Variable length. Maximum length n

SQL (“sequel”) – Cont.

- Provides standard types. Examples:
 - BOOLEAN
 - DATE, TIME, TIMESTAMP
 - DATE: Stores year, month, and day values
 - TIME: Stores hour, minute, and second values
 - TIMESTAMP: Stores year, month, day, hour, minute, and second values
- Additional types differ by vendor:
 - SQLite: <http://www.sqlite.org/datatype3.html>

SQL statements

- create table ...
- drop table ...
- alter table ... add/remove ...
- insert into ... values ...
- delete from ... where ...
- update ... set ... where ...
- select ... from ... where

create table ...

```
CREATE TABLE Company(  
    name VARCHAR(20) PRIMARY KEY,  
    country VARCHAR(20),  
    employees INT,  
    for_profit CHAR(1));
```

Multi-column Keys

- This makes name a key:

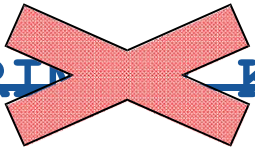
```
CREATE TABLE Company (  
    name VARCHAR(20) PRIMARY KEY,  
    country VARCHAR(20),  
    employees INT,  
    for_profit BOOLEAN);
```

- How can we make a key on name & country?

Multi-column Keys

- Syntax change if a primary key has multiple columns:

```
CREATE TABLE Company (  
  name VARCHAR(20) PRIMARY KEY,  
  country VARCHAR(20),  
  employees INT,  
  for_profit BOOLEAN,  
  PRIMARY KEY (name, country));
```



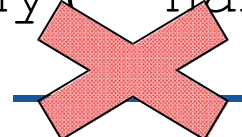
goes away

added

Multi-column Keys (2)

- Likewise for secondary keys:

```
CREATE TABLE Company(  
  name VARCHAR(20),  
  UNIQUE,  
  country VARCHAR(20),  
  employees INT,  
  for_profit BOOLEAN,  
  UNIQUE (name, country));
```




The diagram illustrates the modification of a table's constraints. A large red 'X' is placed over the **UNIQUE** constraint on the **name** column, with a yellow box labeled "goes away" pointing to it. A yellow box labeled "added" points to the **UNIQUE** constraint on the multi-column key **(name, country)**.

Multi-column Keys (3)

- This makes manufacturer a foreign key:

```
CREATE TABLE Product (  
    name VARCHAR(20),  
    price DECIMAL(10,2),  
    manufacturer VARCHAR(20)  
    REFERENCES Company(name) );
```

good idea to include
target column name



Multi-column Keys (3)

- Similar syntax for foreign keys:

```
CREATE TABLE Product (  
    name VARCHAR(20),  
    price DECIMAL(10,2),  
    manu_name VARCHAR(20),  
    manu_co VARCHAR(20),  
    FOREIGN KEY (manu_name, manu_co)  
    REFERENCES Company(name, country) );
```

now need both
name & country

added

UNIQUE

- PRIMARY KEY adds implicit “NOT NULL” constraint while UNIQUE does not
 - you would have to add this explicitly for UNIQUE:

```
CREATE TABLE Company (  
    name VARCHAR(20) NOT NULL, ...  
    UNIQUE (name) );
```

- You almost always want to do this (in real schemas)
 - SQL Server behaves strangely with NULL & UNIQUE
 - otherwise, think through NULL for every query
 - you can remove the NOT NULL constraint later

drop table ...

DROP TABLE Company;

alter table ... add/remove ...

```
ALTER TABLE Company  
ADD CEO VARCHAR(20);
```

insert into ... values ...

```
INSERT INTO Company VALUES  
('GizmoWorks', 'USA', 20000, 'y');
```

One Way to Input Data

- Write a program that outputs SQL statements:

```
for (int a = 1; a <= 50; a++)  
    for (int b = 1; b <= 50; b++)  
        System.out.format(  
            "INSERT INTO T VALUES (%d,%d) ; \n",  
            a, b);
```

- Feed those into SQLite:

```
sqlite3 foo.db < inputs.sql
```

Demo: MakeTriples.java

Warning

- Be very careful when doing this with strings:



```
System.out.format(  
    "INSERT INTO T2 VALUES (%d, '%s') ;",  
    3, "O' Shaughnessy");
```



Becomes:

```
INSERT INTO T2 VALUES (3, 'O' Shaughnessy');
```

which is a syntax error in this case



<https://xkcd.com/327/>

Warning (cont)

- Be very careful when doing this with strings:

```
System.out.format(  
    "INSERT INTO T VALUES (%d, '%s') ;",  
    3, "O' Shaughnessy") ;
```

- This allows a SQL injection attack!
 - Must check for quotes and escape (or disallow) them.
 - We'll see safer ways to do this using JDBC
- DBMSs usually have faster ways to input data
 - SQLite has `.import` (try with `.mode csv`)

delete from ... where ...

```
DELETE FROM Company  
where name = 'GizmoWorks';
```

update ... set ... where ...

```
UPDATE Company  
SET employees = employees + 120  
where name = 'GizmoWorks';
```

select ... from ... where ...

```
SELECT * FROM Company  
where name = 'GizmoWorks';
```

DISTINCT and ORDER BY

- Query results do not have to be relations
 - i.e., they can have duplicate rows
 - remove them using DISTINCT
- Result order is normally unspecified
 - choose an order using ORDER BY
 - e.g., ORDER BY country, cname
 - e.g., ORDER BY price ASC, pname DESC
- Examples in `lec03-sql-basics.sql`

Demo on Sqlite

- E.g., type `sqlite3` in Cygwin
- `.exit` - exit from `sqlite3`

SQLite Uses

- SQLite is just a library
- Can be used as part of any C/C++/Java program
 - ex: could be used in an iPhone app
- Can be used in Chrome & Safari
 - no support in Firefox or IE

Physical Data Independence

- SQL doesn't specify how data is stored on disk
- No need to think about encodings of data types
 - ex: DECIMAL(10,2)
 - ex: VARCHAR(255)
 - does this need to use 255 bytes to store 'hello'?
- No need to think about how tuples are arranged
 - ex: could be row- or column-major ordered
 - (Most DBMSs are row-ordered, but Google's BigQuery is column-oriented.)

SQLite Gotchas

- Allows NULL keys
 - At most one tuple can have NULL in the key
 - According to the SQL standard, PRIMARY KEY should always imply NOT NULL, but this is not the case in SQLite
- Does not support boolean or date/time columns
- Doesn't always enforce domain constraints!
 - will let you insert a string where an INT is expected
- Doesn't enforce foreign key constraints by default
- Etc...