

# Introduction to Database Systems

Lecture 09: E/R Diagrams (4.1-6)  
and Constraints (7.1-2)

# Database Design

What it is:

- Starting from scratch, design the database schema: relation, attributes, keys, foreign keys, constraints etc.

Why it's hard:


- The database will be in operation for years.
- Updating the schema in production is very hard:
  - schema change modifications are expensive (why?)
  - making the change without introducing any bugs is hard
    - this part is, by far, the most important consideration in practice

# Database Design


- Consider issues such as:
  - What entities to model
  - How entities are related
  - What constraints exist in the domain
- Several formalisms exists
  - We discuss E/R diagrams
- Reading: Sec. 4.1-4.6

# Database Design Process

Conceptual Model:




Relational Model:  
Tables + constraints  
And also functional dep.




Normalization:  
Eliminates anomalies

Conceptual Schema



Physical storage details

Physical Schema



# Entity / Relationship Diagrams

- Entity set = a class
  - An entity = an object



Product

- Attribute

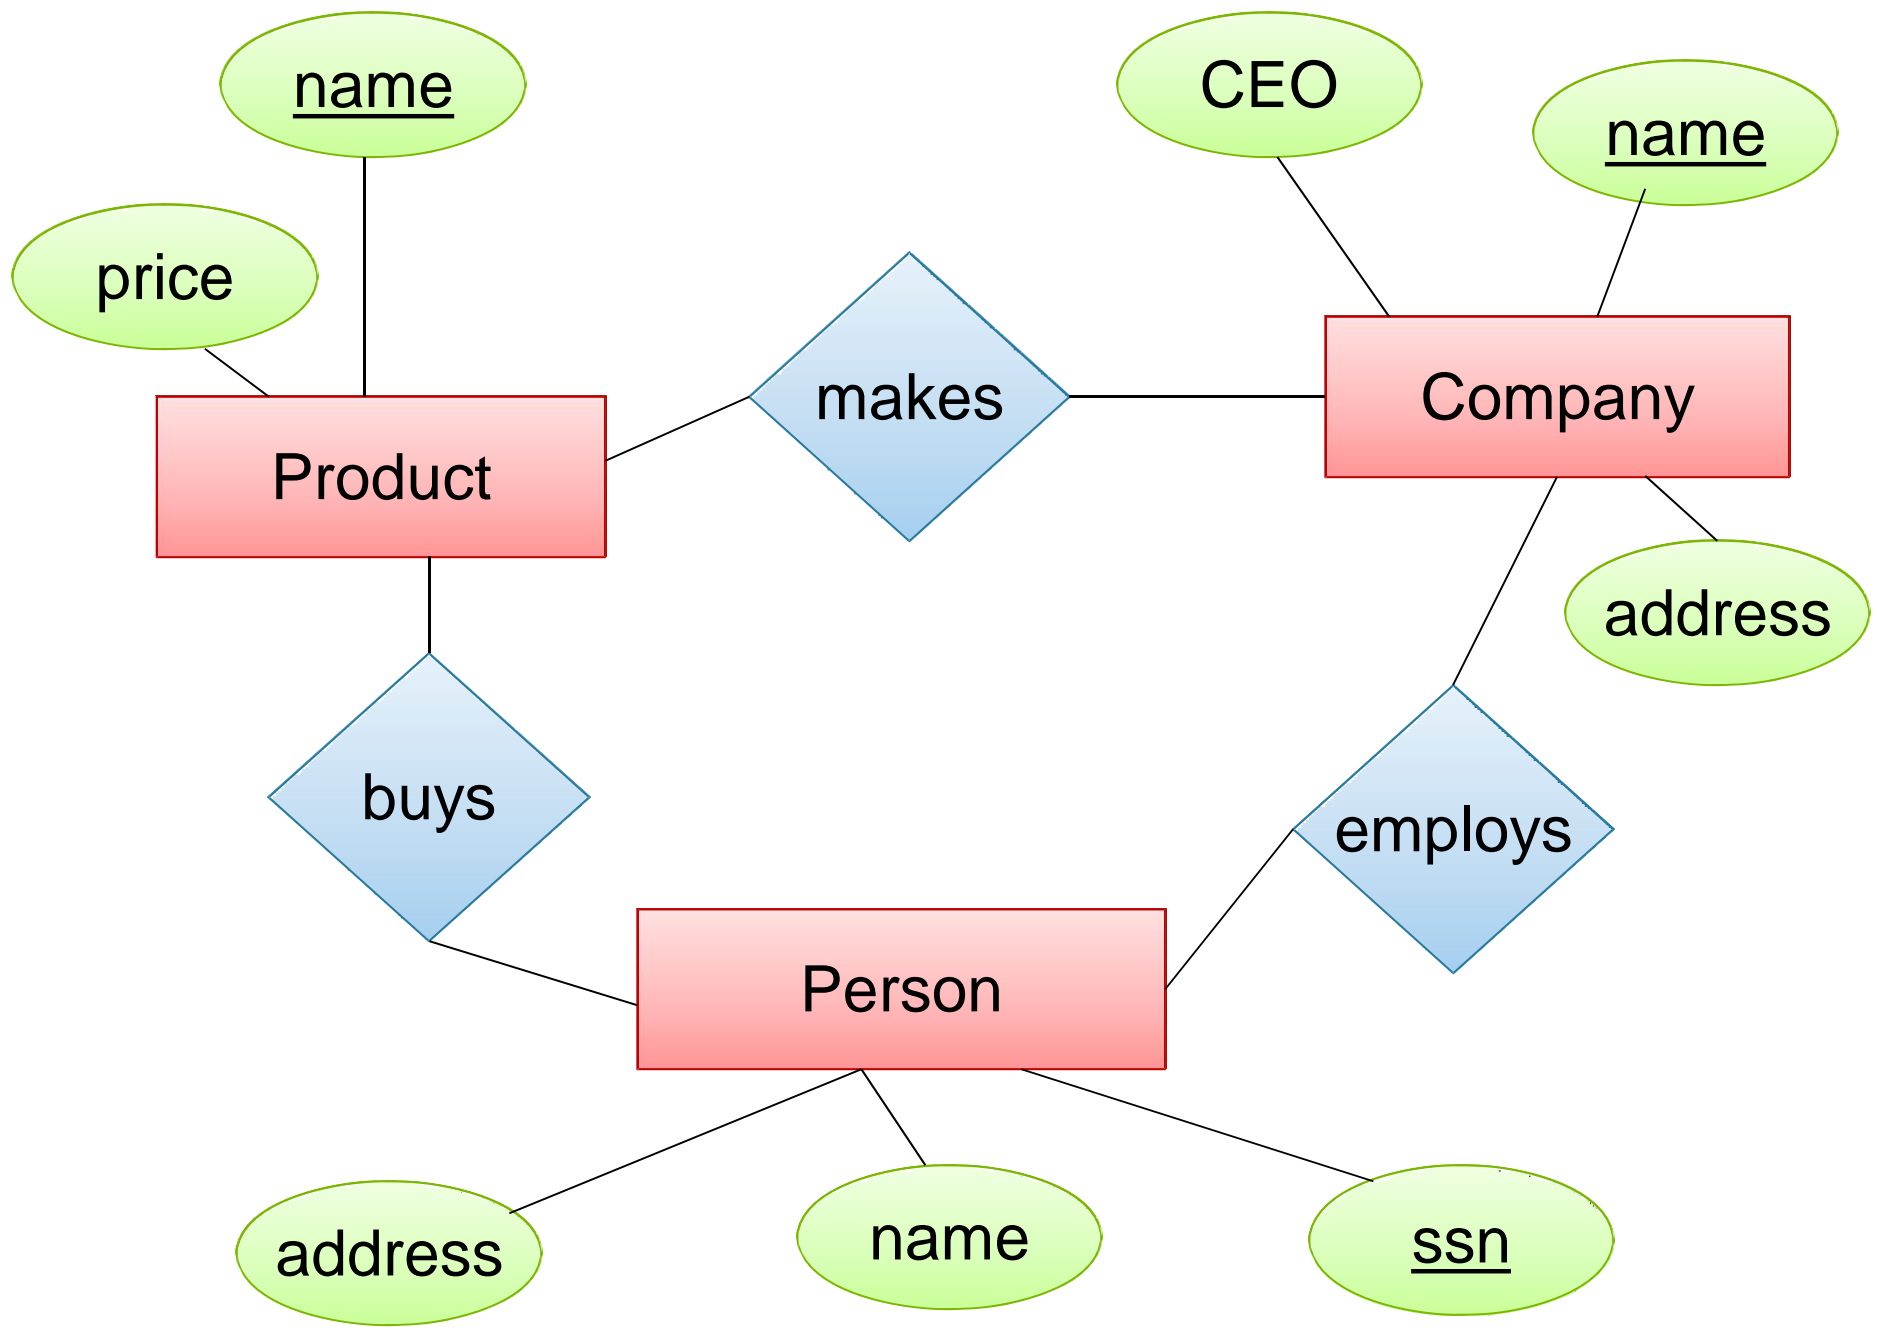


city

- Relationship




makes



# Keys in E/R Diagrams


- Every entity set must have a key




# What is a Relation(ship)?

- A mathematical definition:
  - if  $A, B$  are sets, then a relation  $R$  is a subset of  $A \times B$

- $A = \{1, 2, 3\}$ ,  $B = \{a, b, c, d\}$ ,  
 $A \times B = \{(1,a), (1,b), \dots, (3,d)\}$   
 $R = \{(1,a), (1,c), (3,b)\}$




- **makes** is a subset of **Product**  $\times$  **Company**:






# Multiplicity of E/R Relations





- one-one:

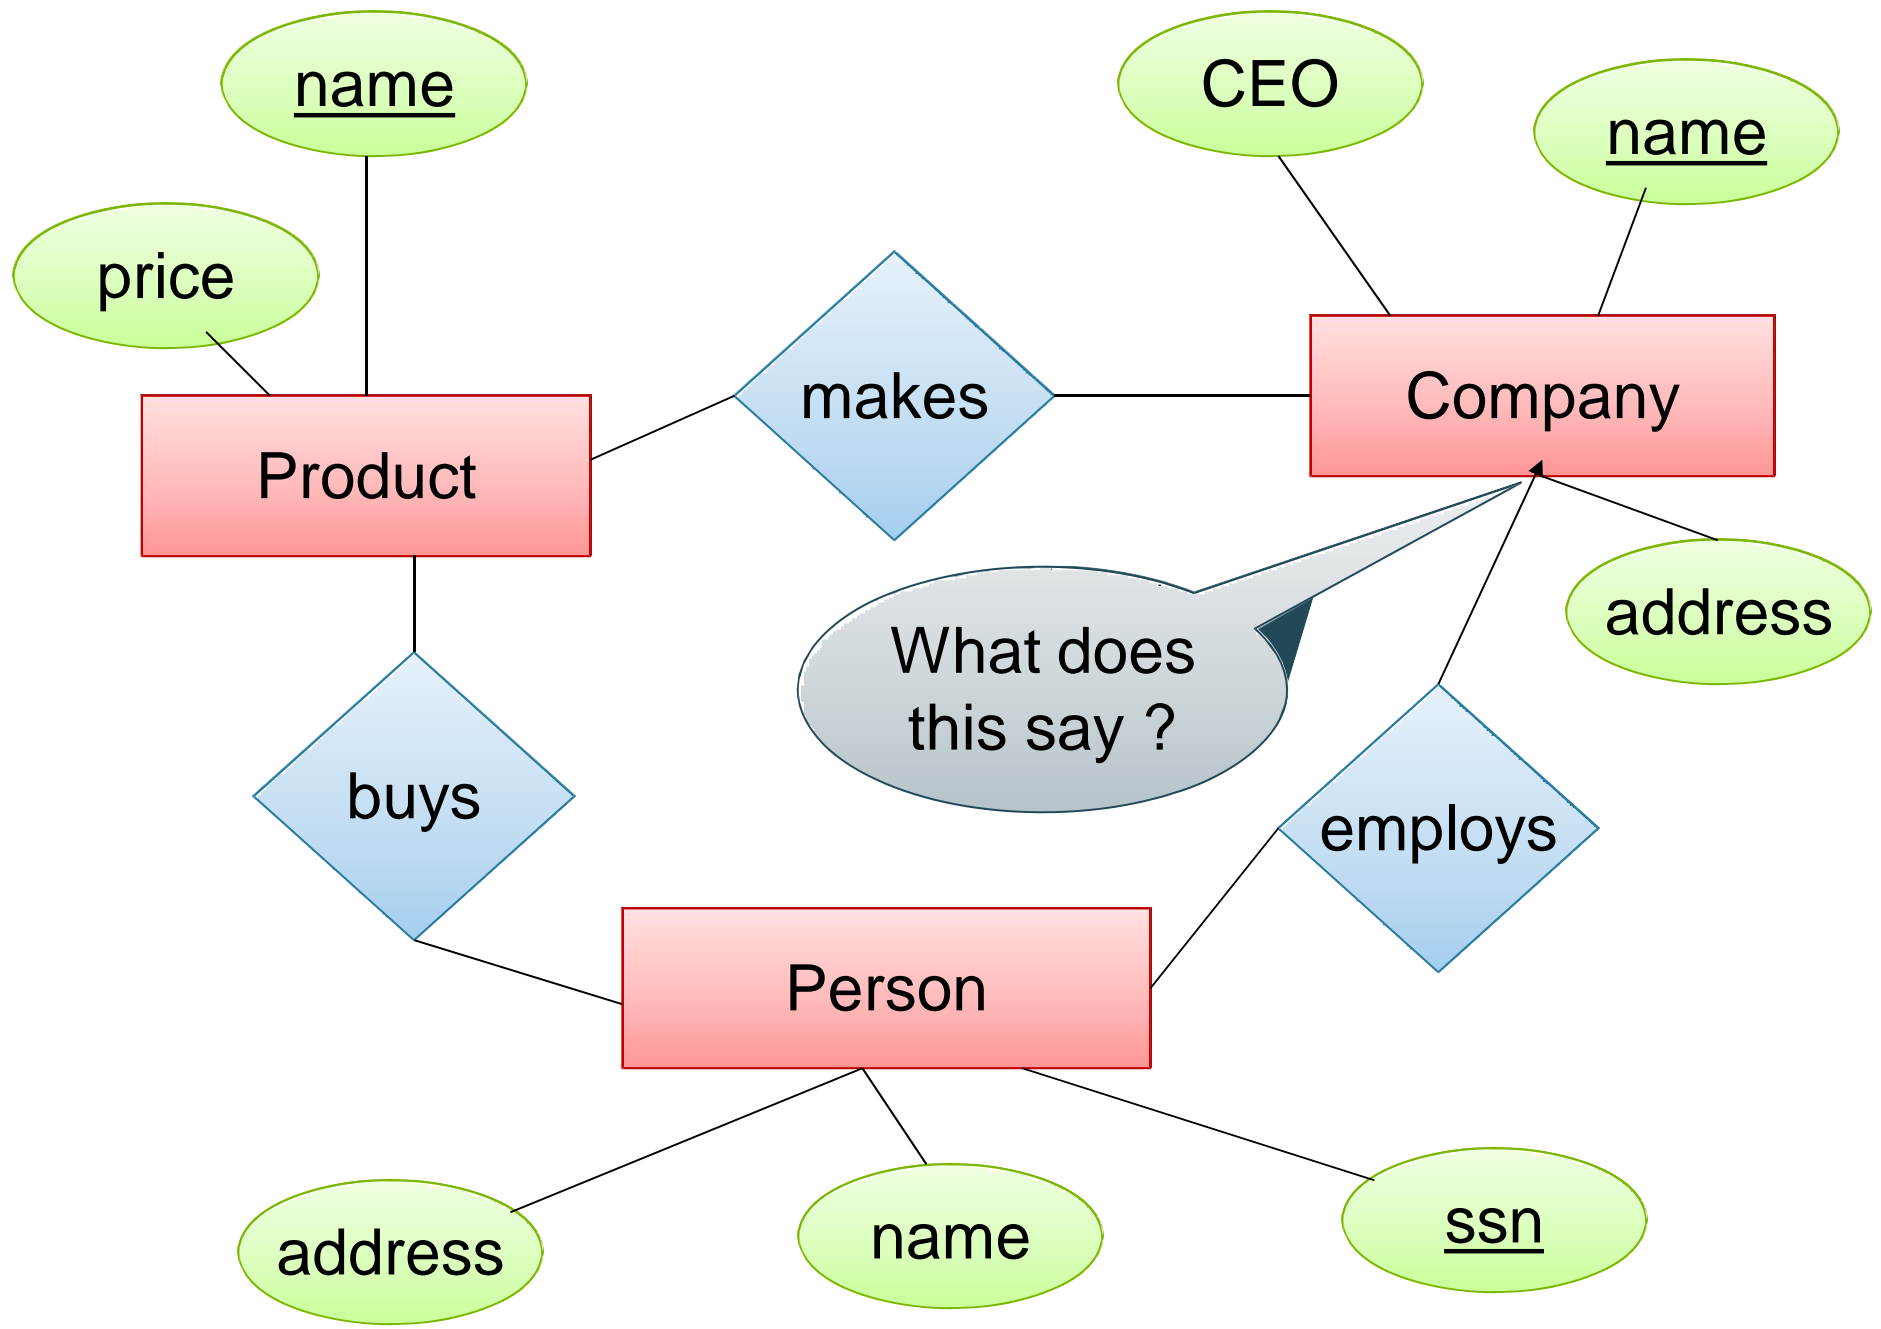


- many-one



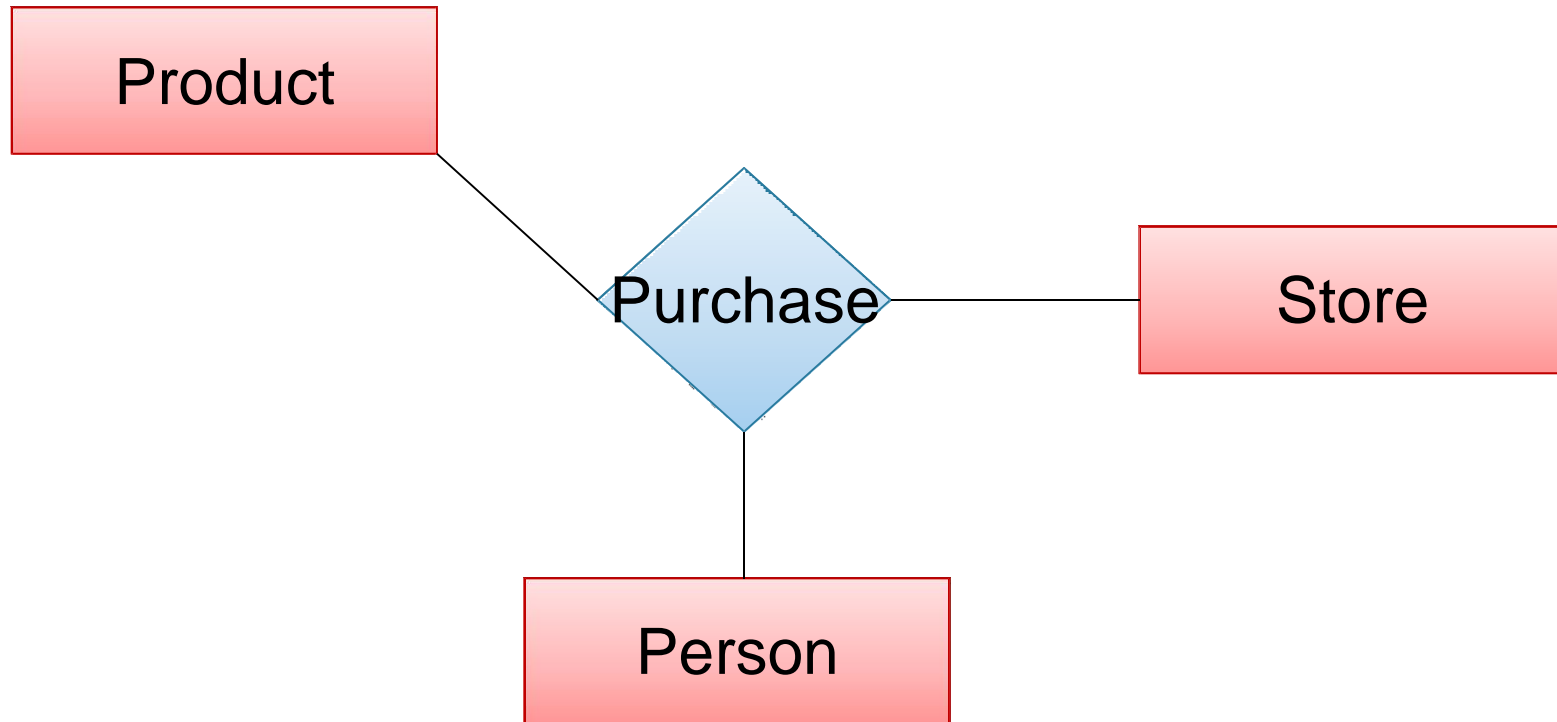
- many-many





# Multi-way Relationships

How do we model a purchase relationship between buyers, products and stores?




Can still model as a mathematical set (Q. how ?)

A. As a set of triples  $\subseteq \text{Person} \times \text{Product} \times \text{Store}$

# Arrows in Multiway Relationships


**Q:** What does the arrow mean ?



**A:** A given person buys a given product from at most one store  
[Arrow pointing to E means that if we select one entity from each of the other entity sets in the relationship, those entities are related to at most one entity in E]

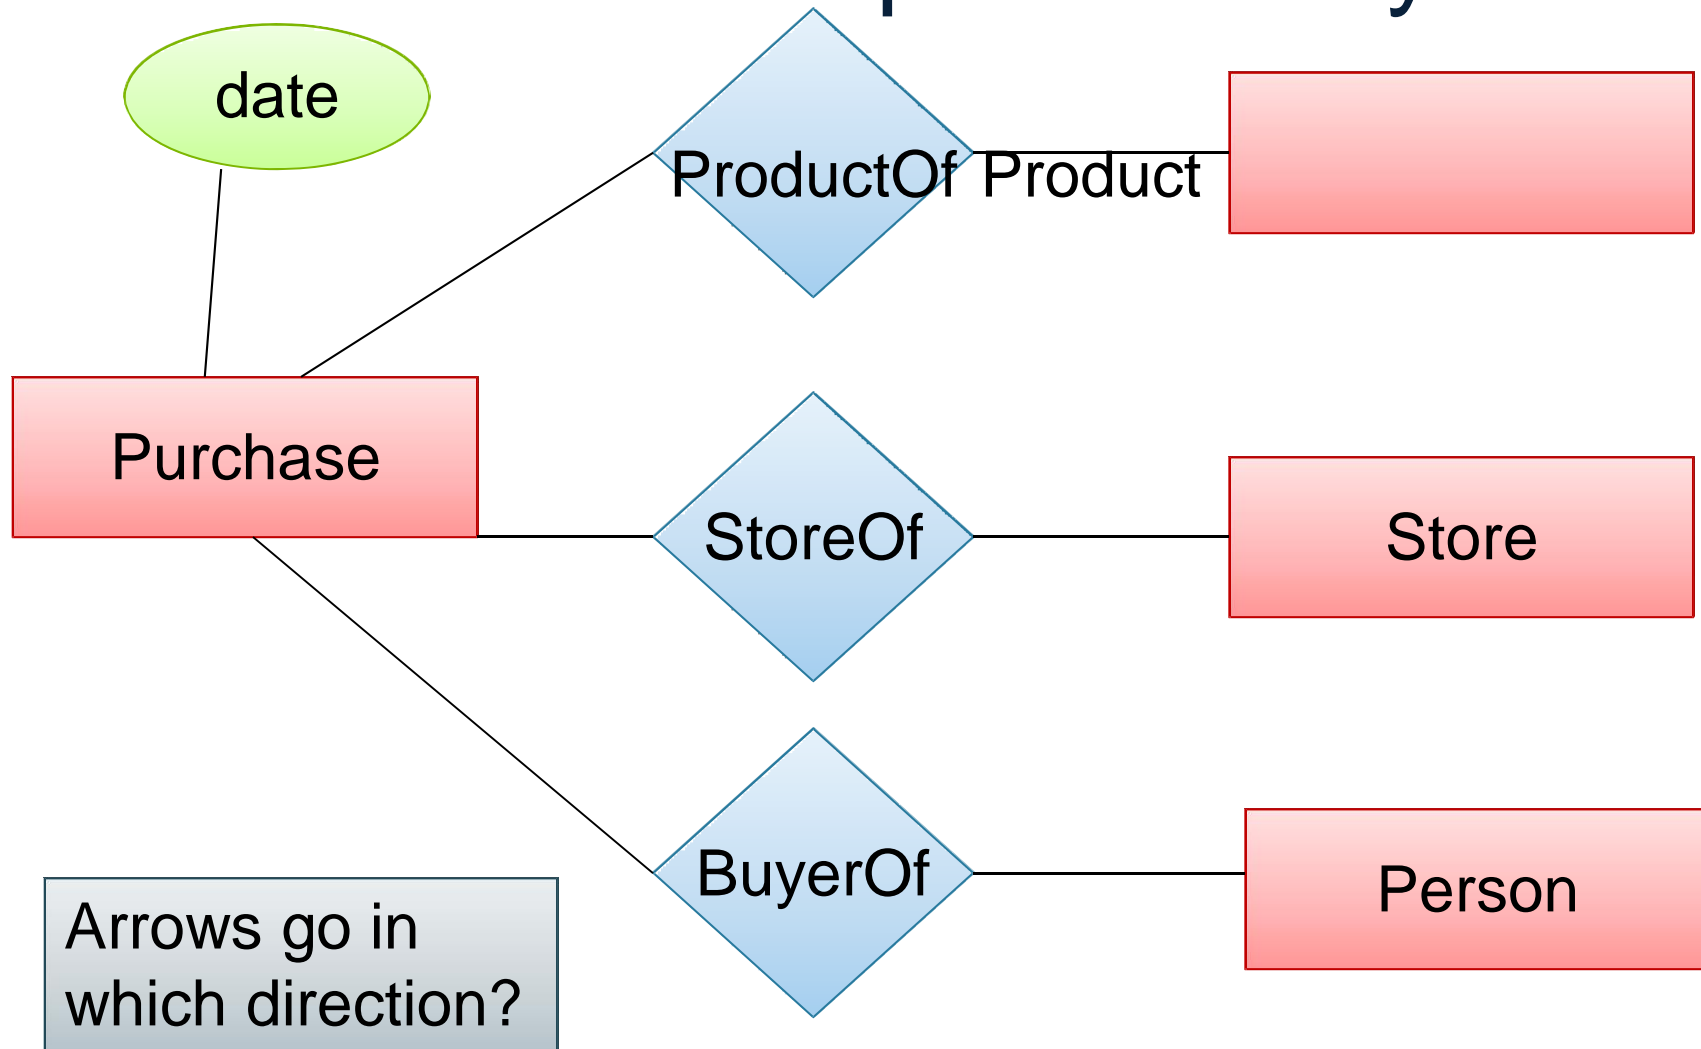
# Arrows in Multiway Relationships

**Q:** What does the arrow mean ?

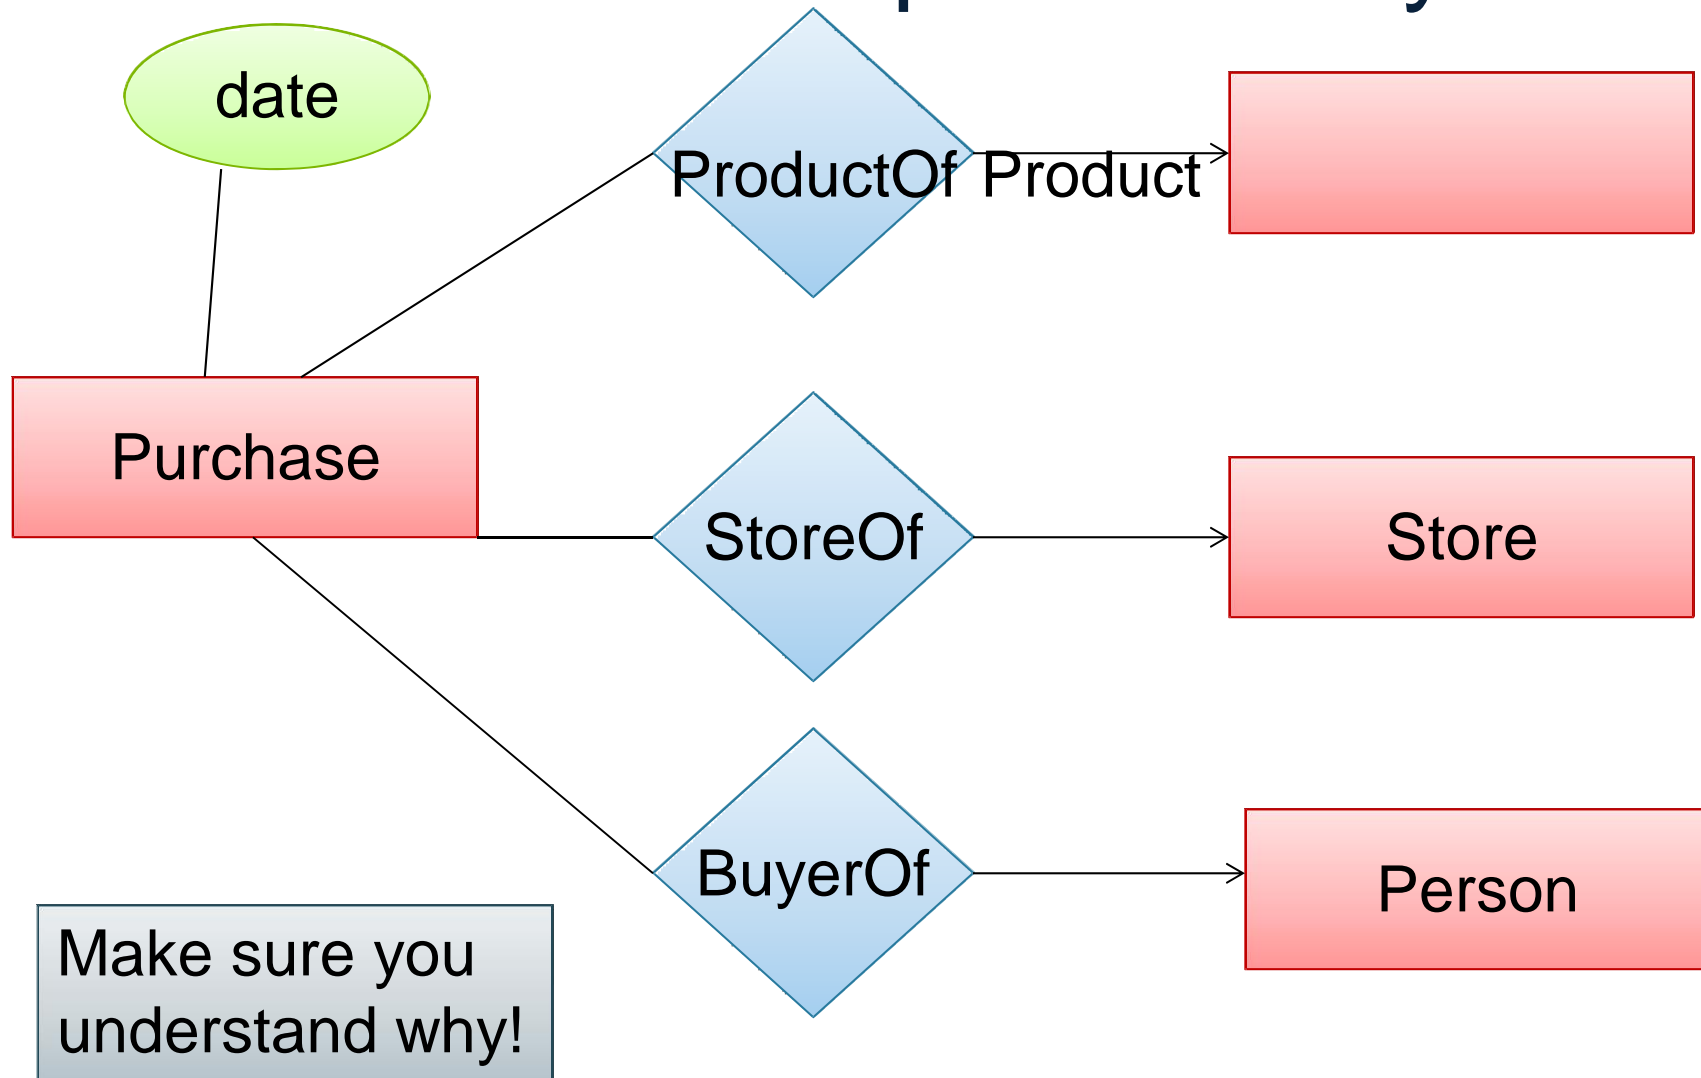


**A:** A given person buys a given product from at most one store AND every store sells to every person at most one product

# Converting Multi-way Relationships to Binary




# Converting Multi-way Relationships to Binary



### 3. Design Principles

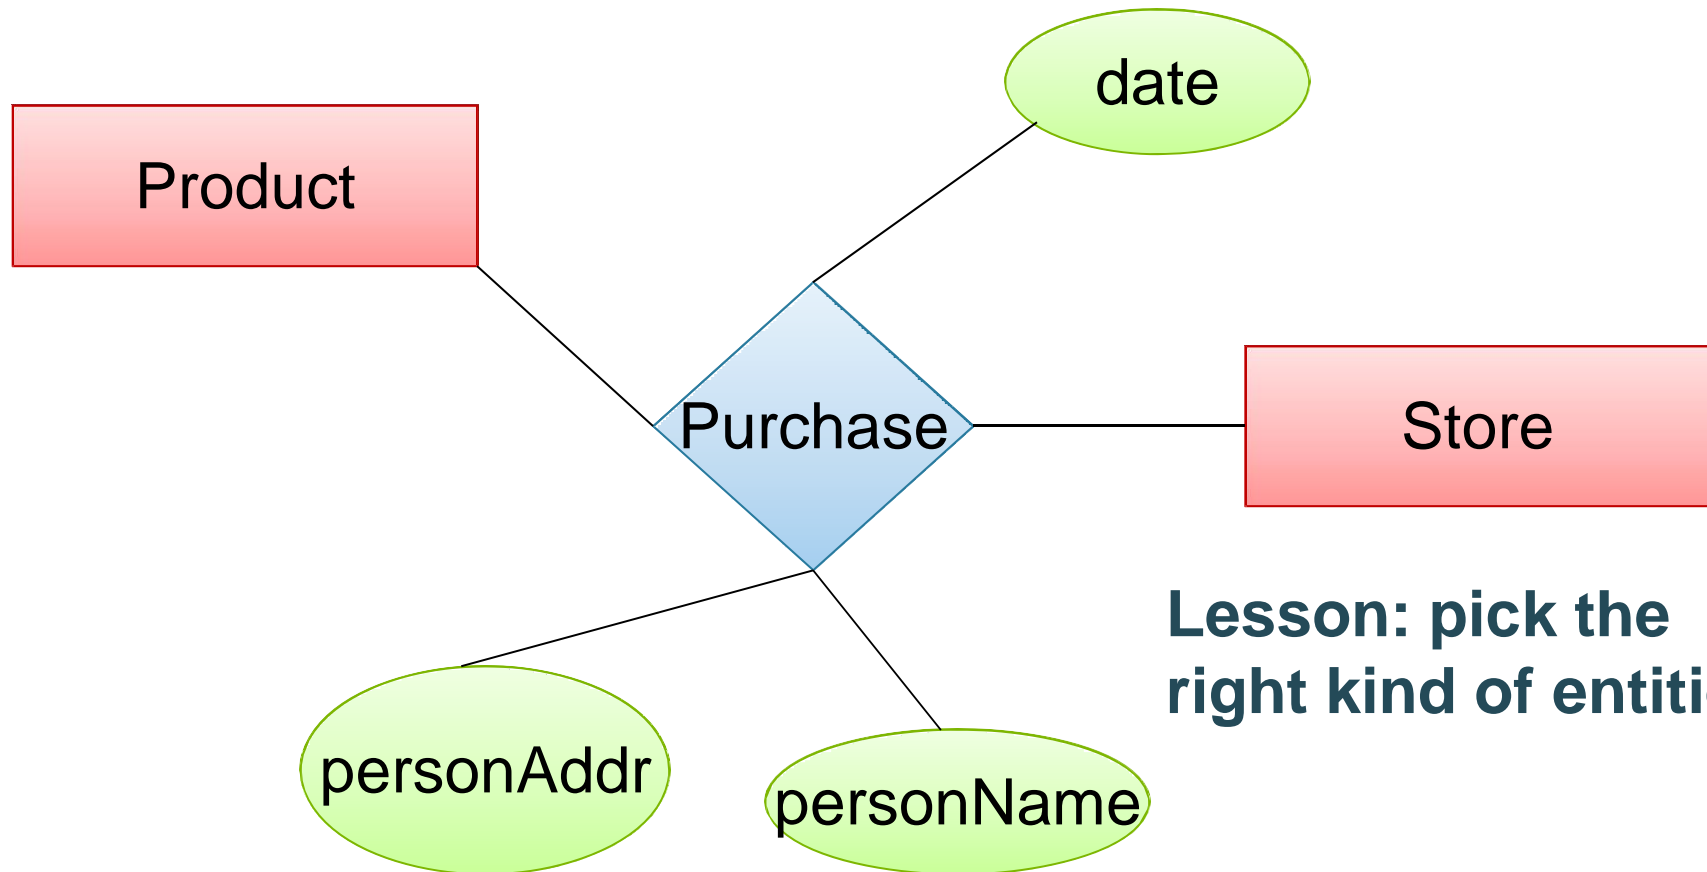
**What's wrong?**



**Lesson: be faithful to the specifications of the app!**

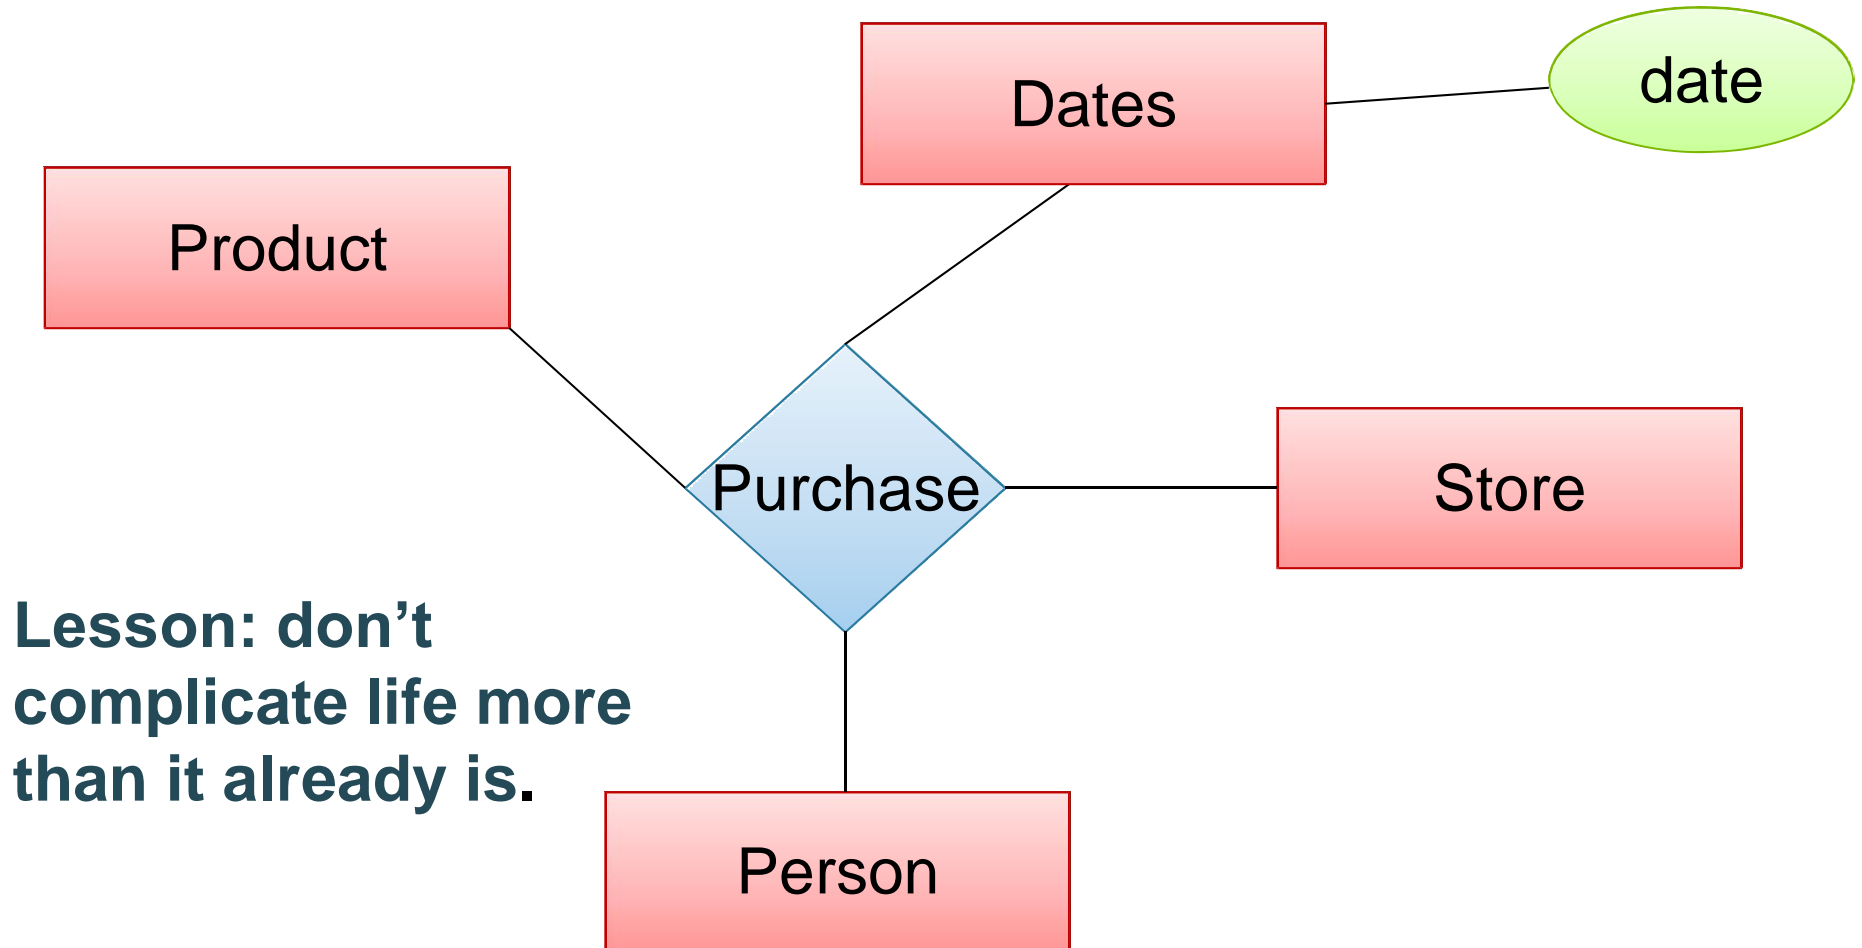


# Design Principles: What's Wrong?



**Lesson: pick the  
right kind of entities.**

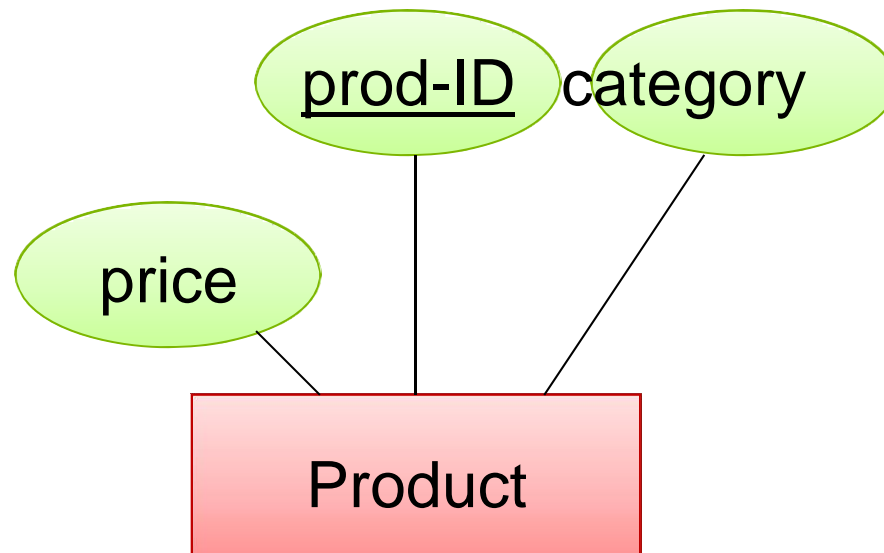
# Design Principles: What's Wrong?



# From E/R Diagrams to Relational Schema

- Entity set  $\rightarrow$  relation
- Relationship  $\rightarrow$  relation

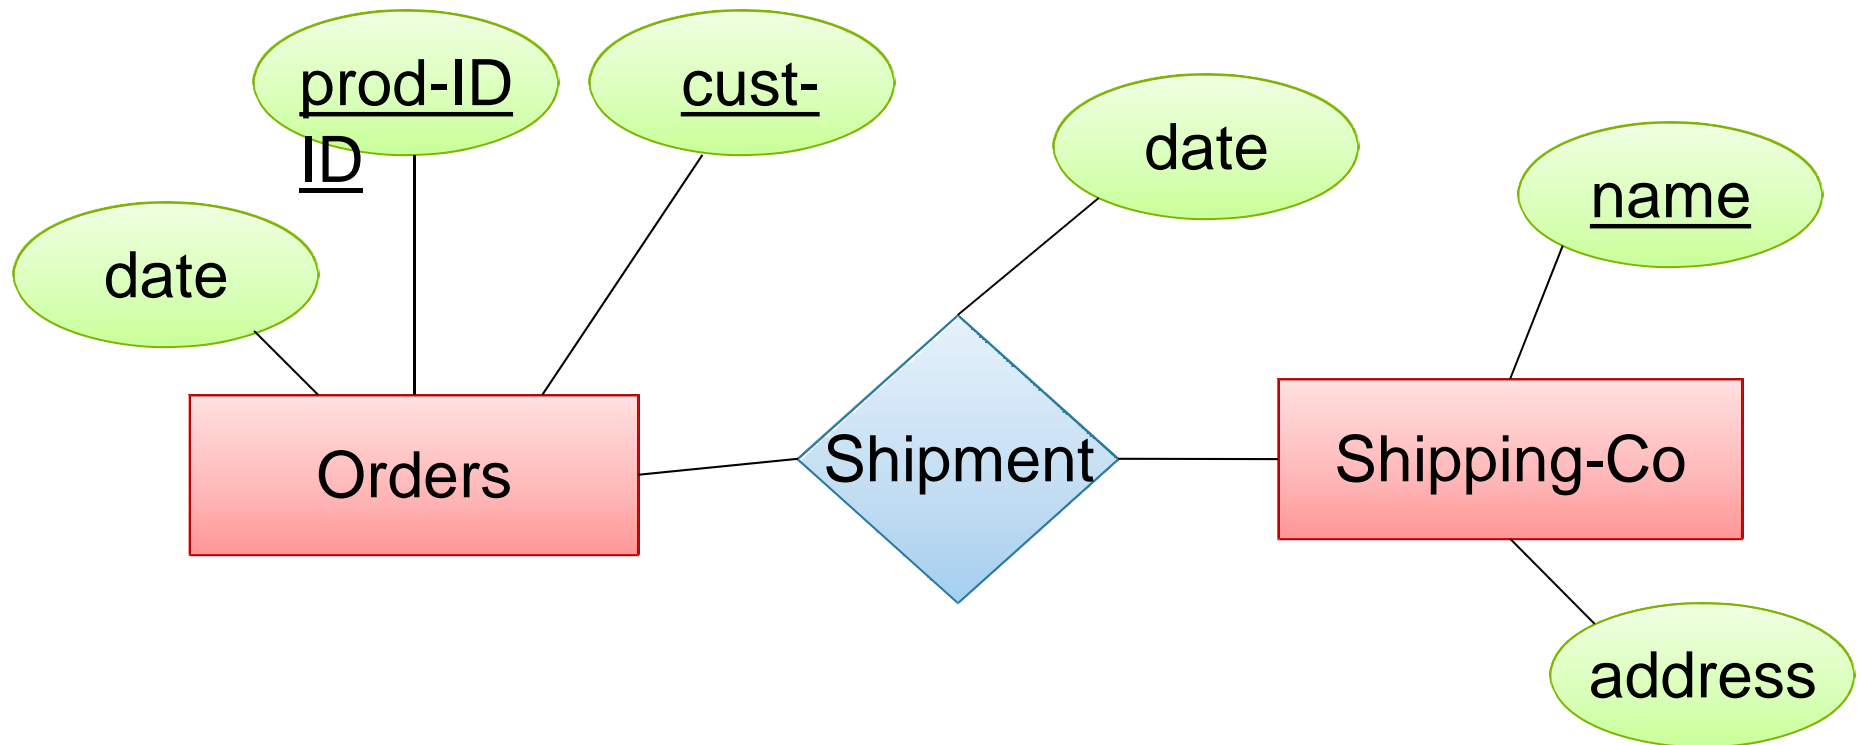
# Entity Set to Relation



**Product**(prod-ID, category, price)

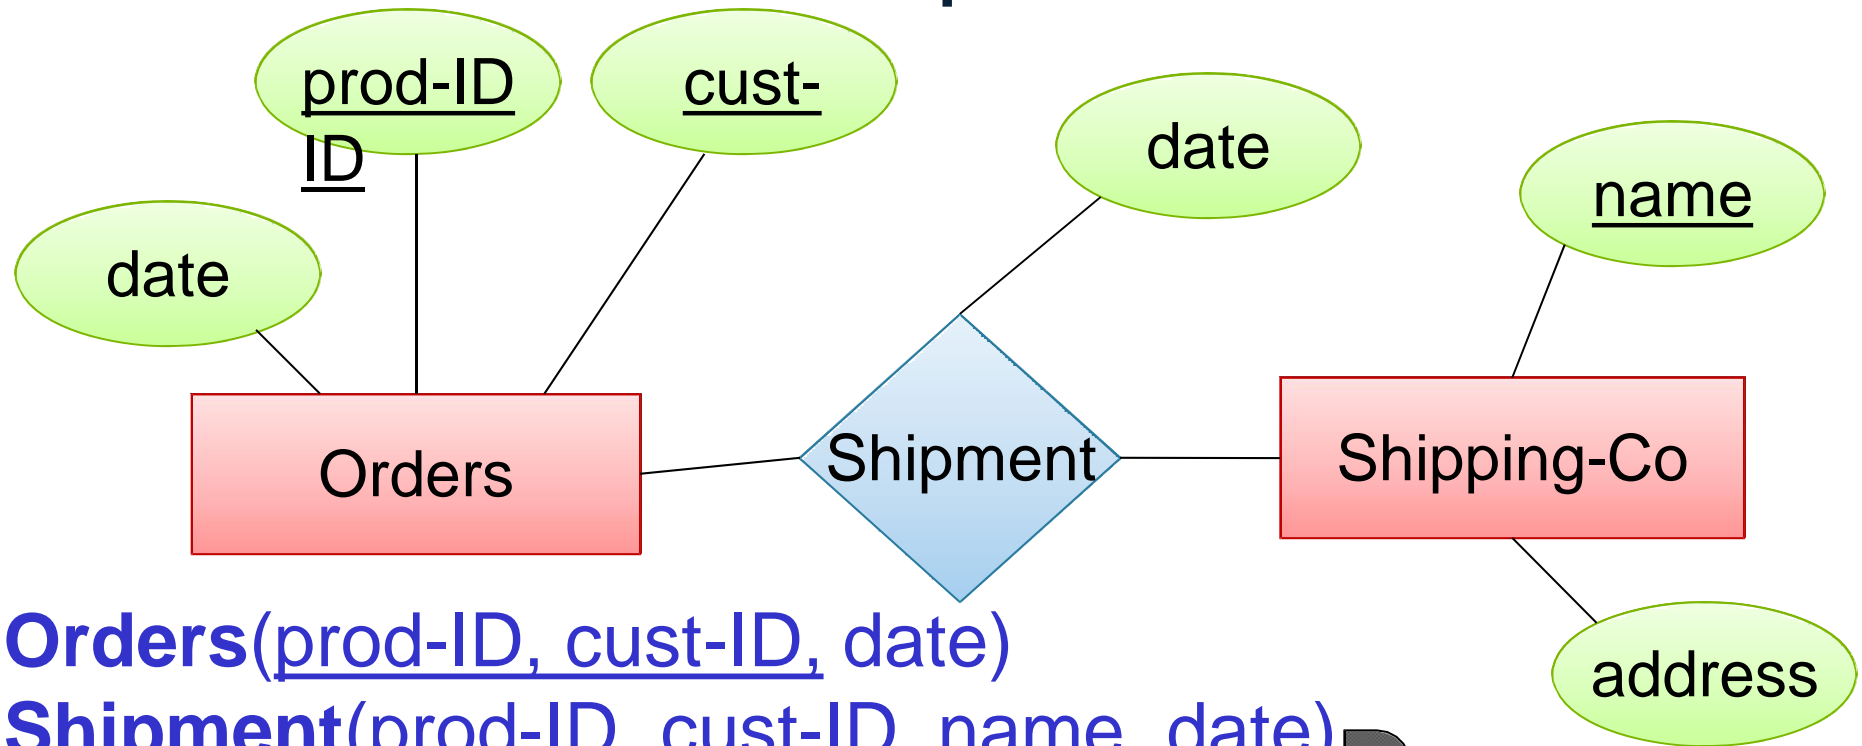
<u>prod-ID</u>	category	price
Gizmo55	Camera	99.99
Pokemn19	Toy	29.99

# N-N Relationships to Relations



Represent this in relations

# N-N Relationships to Relations



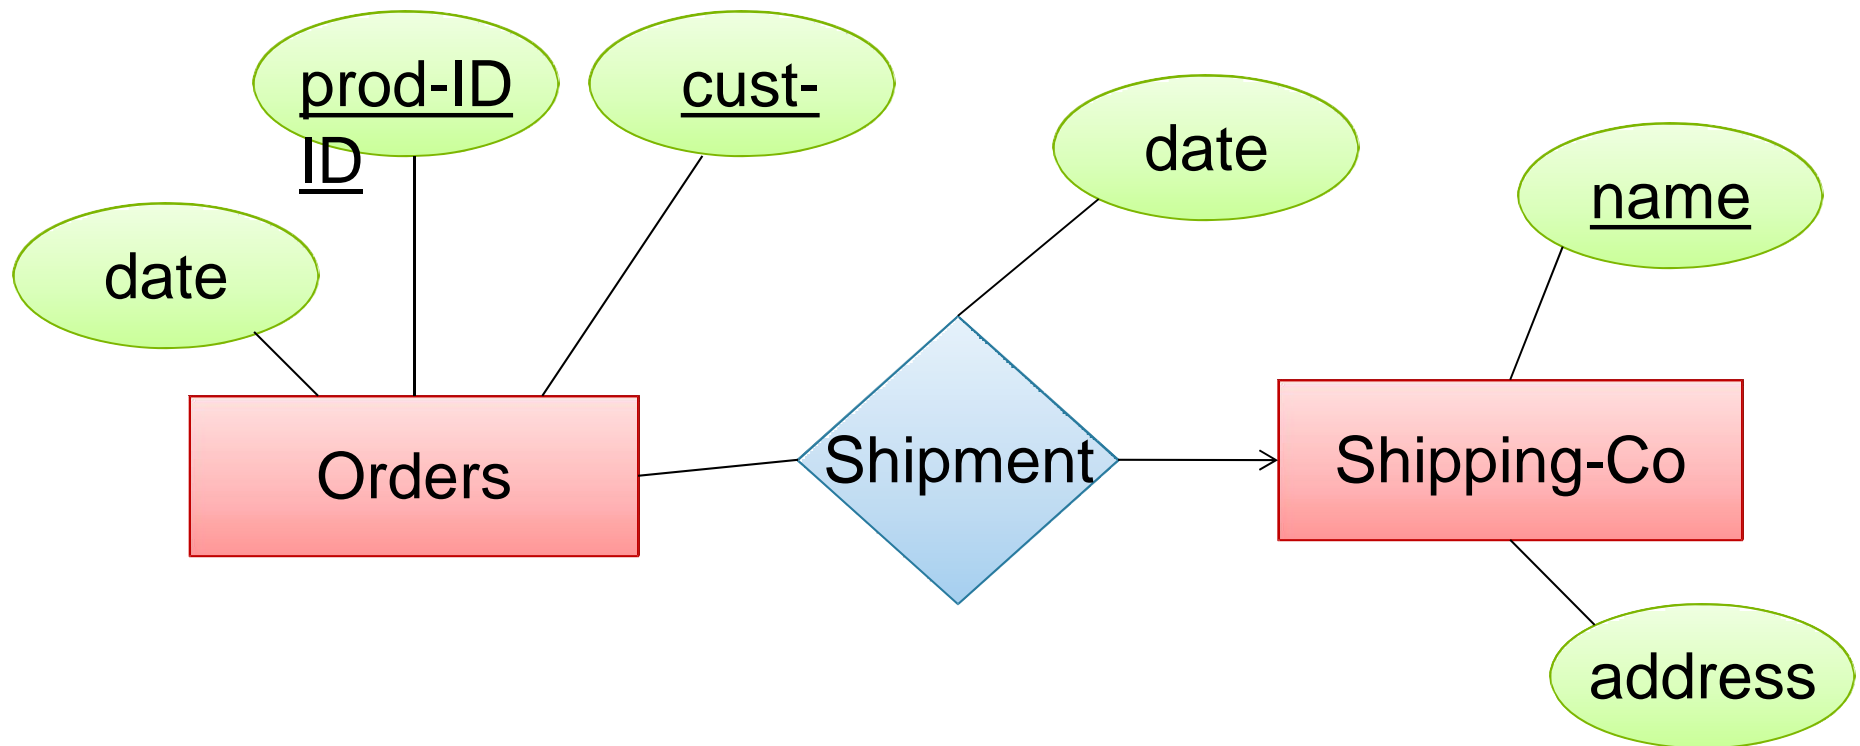
**Orders**(prod-ID, cust-ID, date)

**Shipment**(prod-ID, cust-ID, name, date)

**Shipping-Co**(name, address)

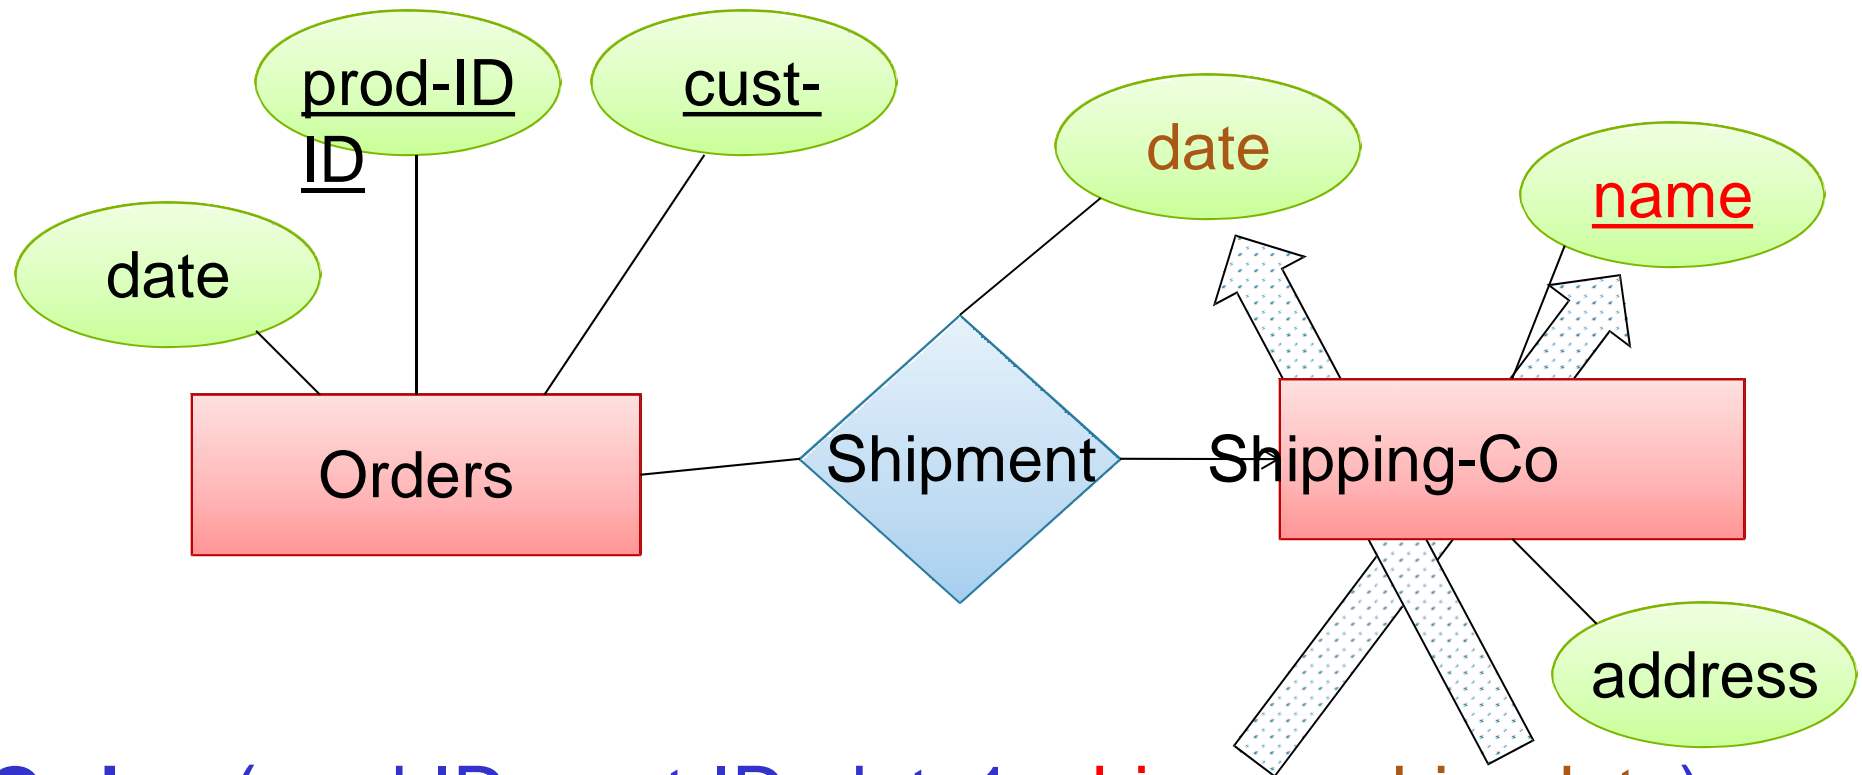
<u>prod-ID</u>	<u>cust-ID</u>	<u>name</u>	<u>date</u>
Gizmo55	Joe12	UPS	4/10/2011
Gizmo55	Joe12	FEDEX	4/9/2011

# N-1 Relationships to Relations



Represent this in relations

# N-1 Relationships to Relations

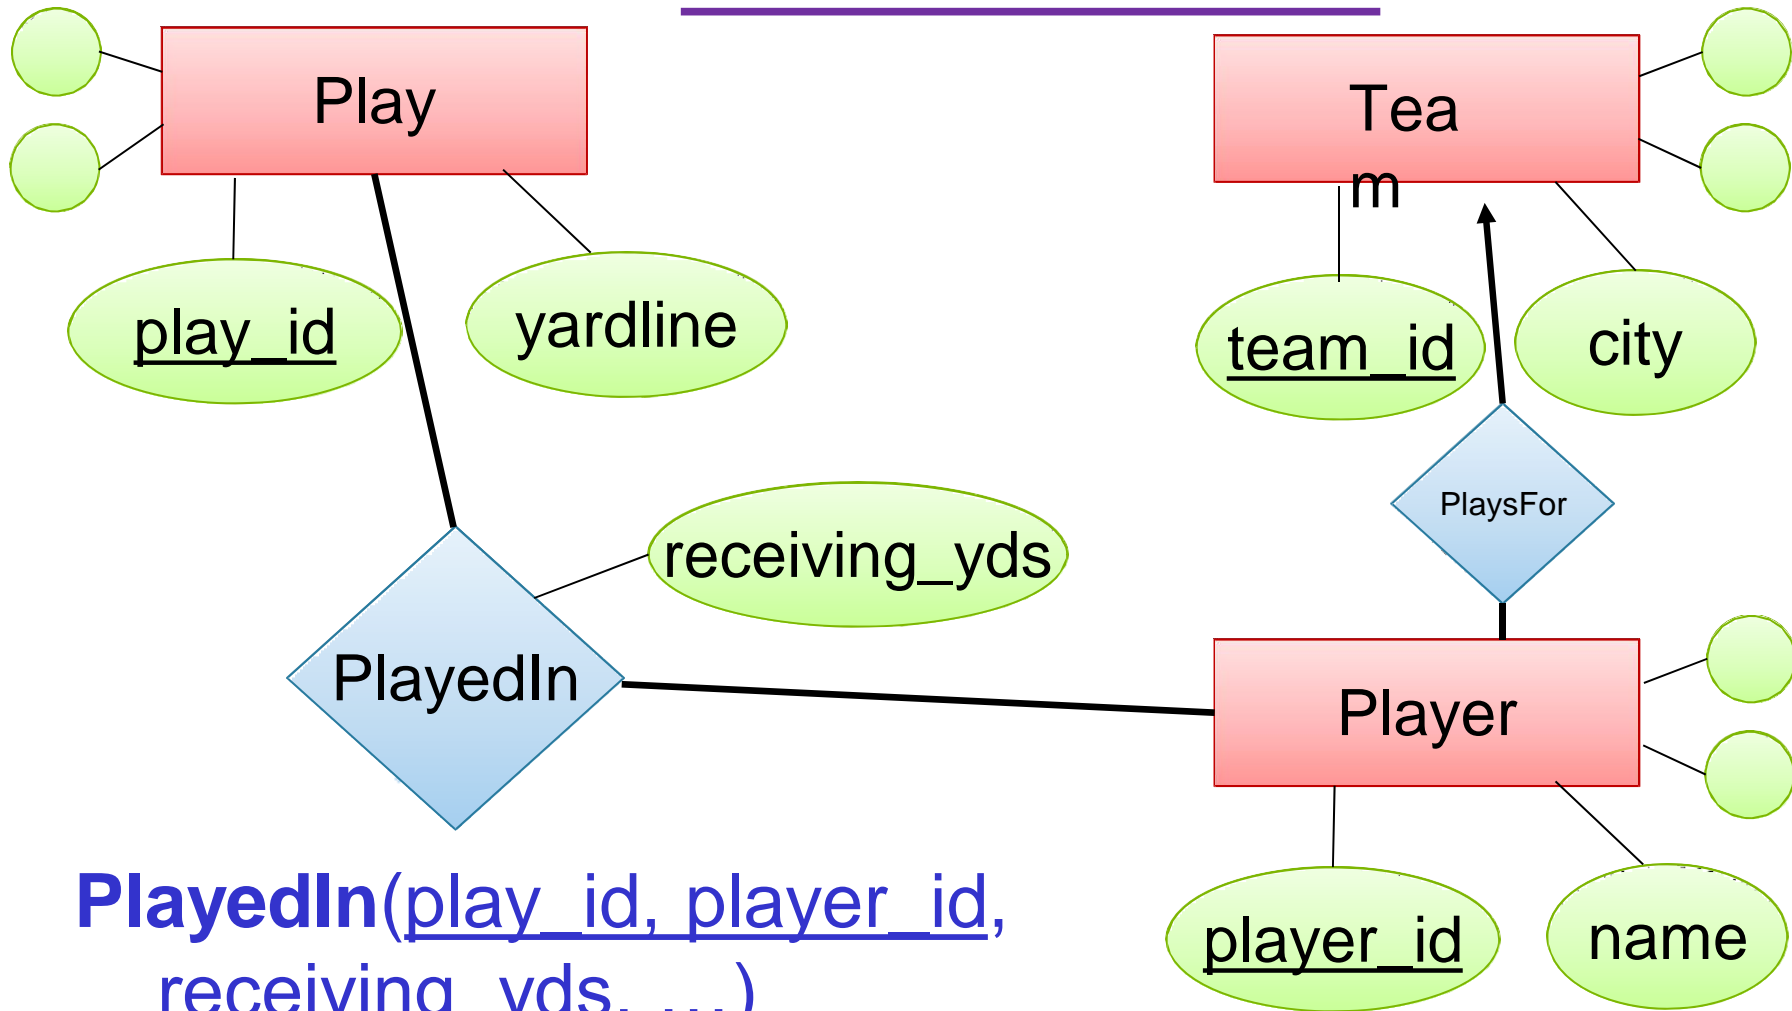


**Orders**(prod-ID, cust-ID, date1, **ship\_co**, ship\_date)  
**Shipping-Co**(name, address)

Remember: many-one relationship becomes FK, not relation



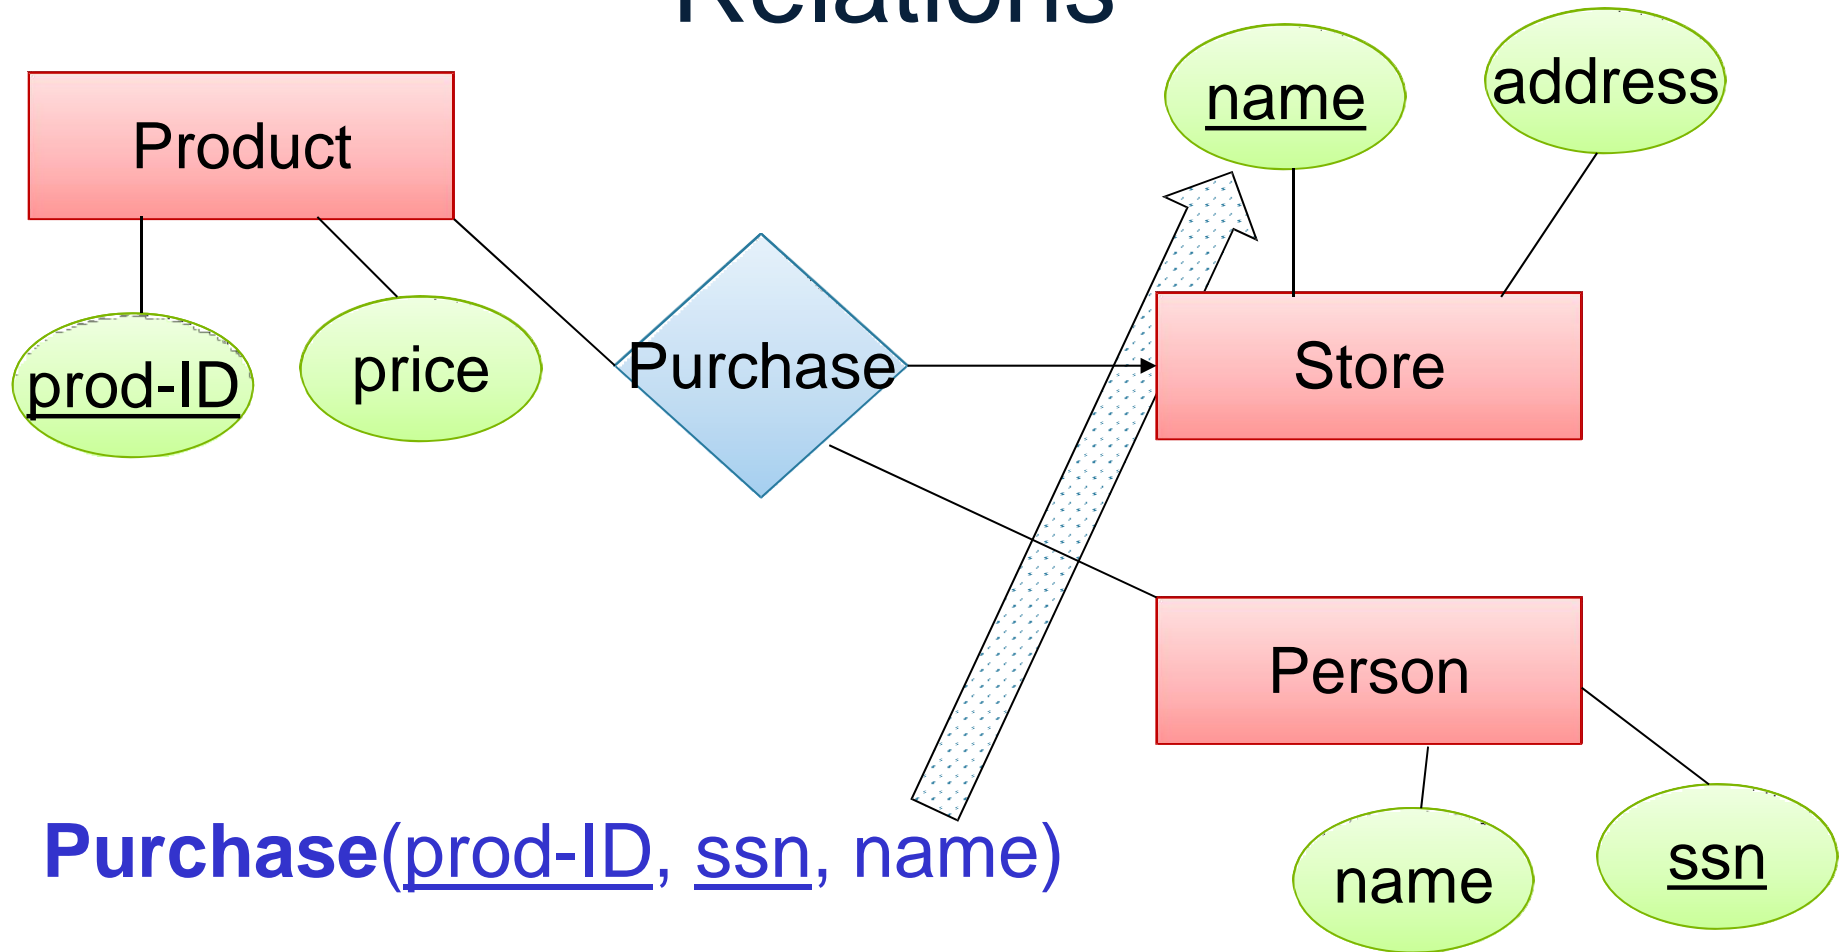
## Ex: NFL Game DB



**PlayedIn(play\_id, player\_id,**  
**receiving\_yds, ...)**

(Actually, the key of Play is not play\_id. More on this later...)


# Multi-way Relationships to Relations



# Modeling Subclasses

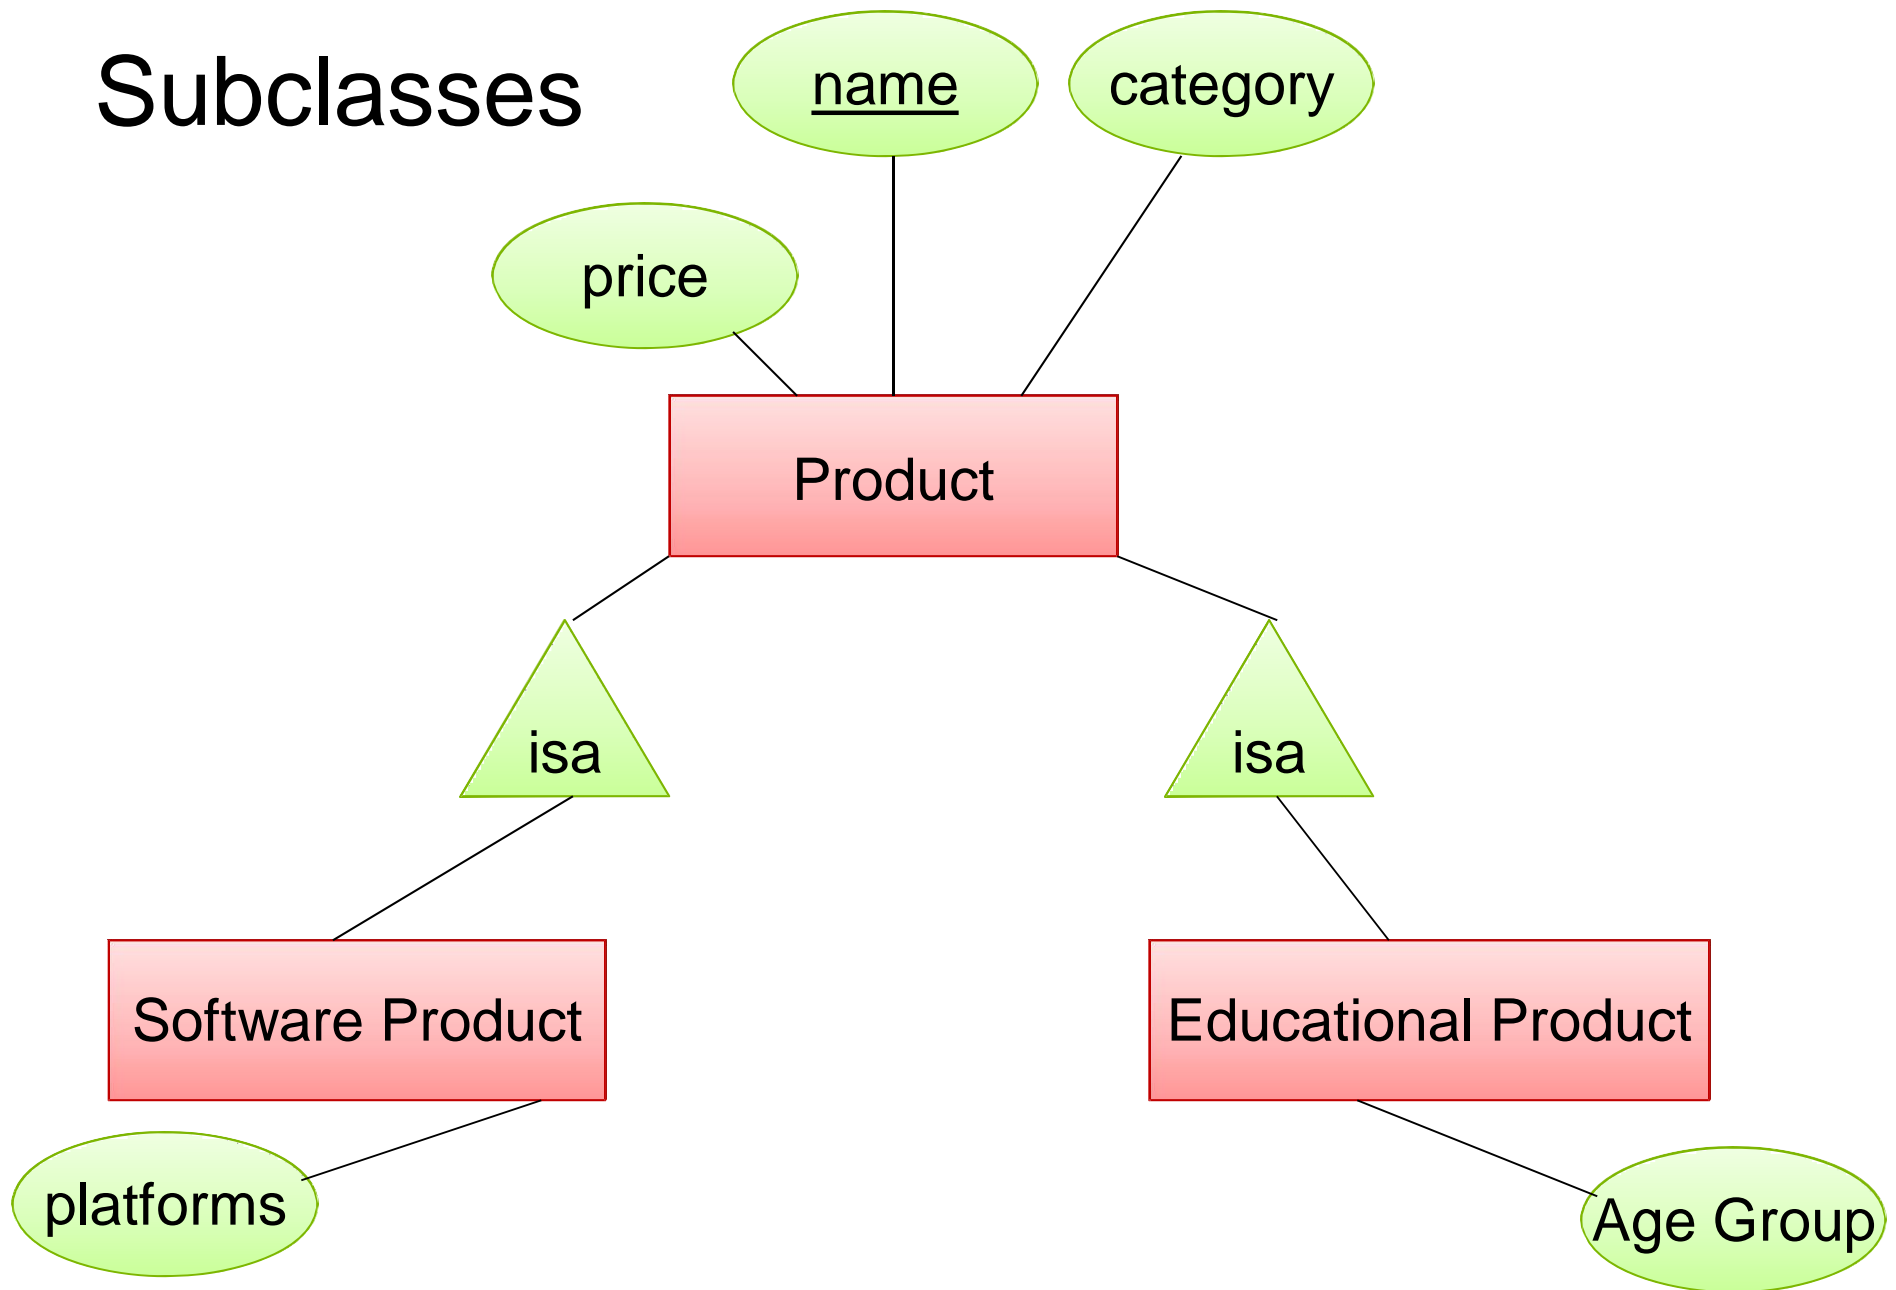
Some objects in a class may be special

- define a new class
- better: define a *subclass*

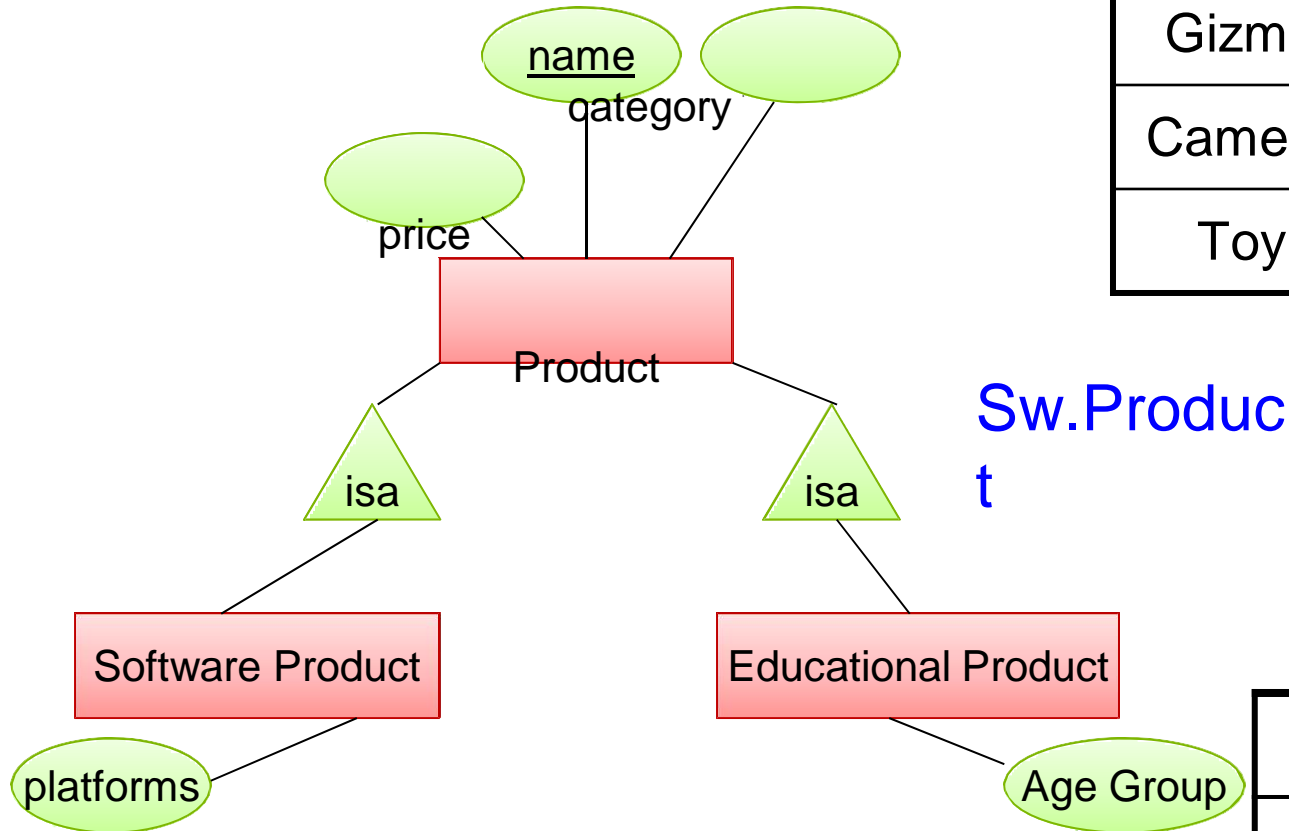


So --- we define subclasses in E/R

# Subclasses



# Subclasses to Relations (one option)



Product

<u>Name</u>	Price	Category
Gizmo	99	gadget
Camera	49	photo
Toy	39	gadget

Sw.Product


<u>Name</u>	platforms
Gizmo	unix

Ed.Product

<u>Name</u>	Age Group
Gizmo	toddler
Toy	retired

Other ways to convert are possible...

# Modeling Union Types with Subclasses




Say: each piece of furniture is owned either by a person or by a company

# Modeling Union Types with Subclasses


Say: each piece of furniture is owned either by a person or by a company

Solution 1. Acceptable but imperfect (What's wrong?)



# Modeling Union Types with Subclasses


Solution 2: better, more laborious





# Weak Entity Sets

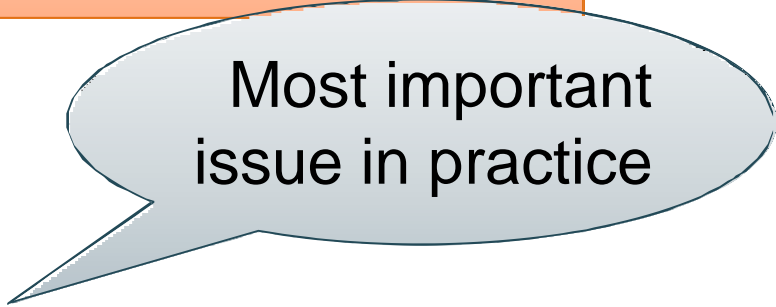
Entity sets are weak when their key comes from other classes to which they are related.



Team(sport, number,  
universityName) University(name)

# Integrity Constraints Motivation

An integrity constraint is a condition specified on a database schema that restricts the data that can be stored in an instance of the database.



Most important  
issue in practice

- ICs help prevent entry of incorrect information
- How? DBMS enforces integrity constraints
  - Allows only legal database instances (i.e., those that satisfy all constraints) to exist
  - Ensures that all necessary checks are always performed and avoids duplicating the verification logic in each application

# Constraints in E/R Diagrams

Finding constraints is part of the modeling process.  
Commonly used constraints:

**Keys:** social security number uniquely identifies a person.

**Single-value constraints:** can have only one genetic father

**Referential integrity constraints:** if you work for a company,

it

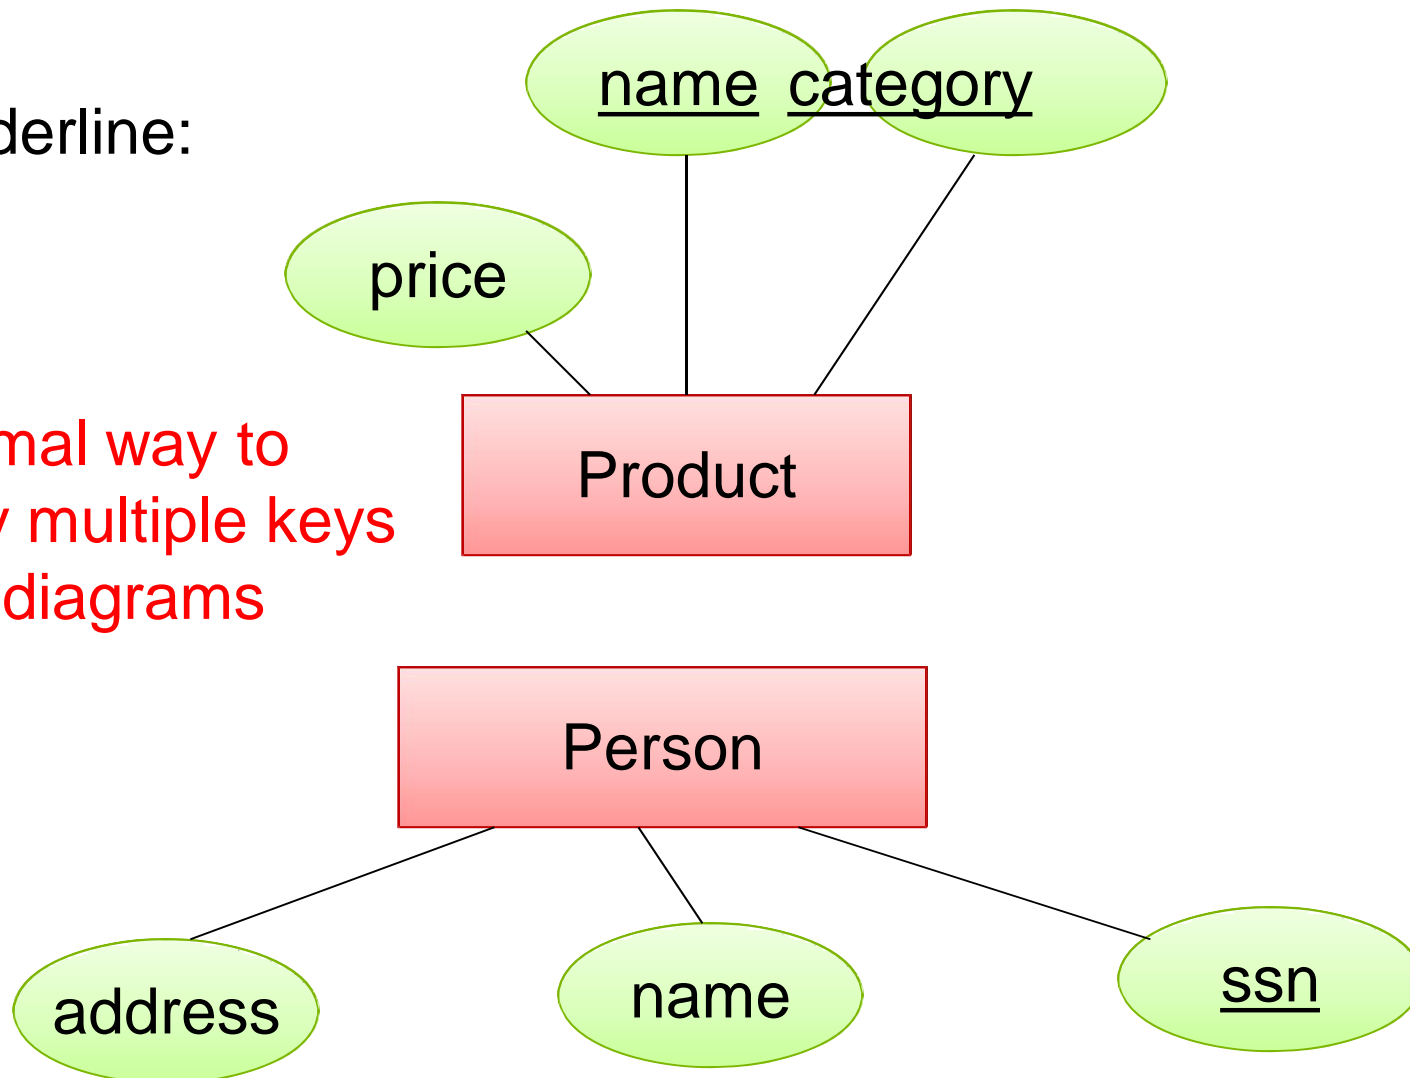
must exist in the database.

**Other constraints:** peoples' ages are between 0 and 150<sub>35</sub>  
some values should not be NULL

# Keys in E/R Diagrams

Underline:

No formal way to  
specify multiple keys  
in E/R diagrams




# Single Value Constraints



vs.




# Referential Integrity Constraints




Each product made by at most one  
Company  
Some products made by no company

Which one is FK?



Each product made by exactly one  
company.

# Other Constraints



Q: What does this mean ?

A: A Company entity cannot be connected by relationship to more than 99 Product entities

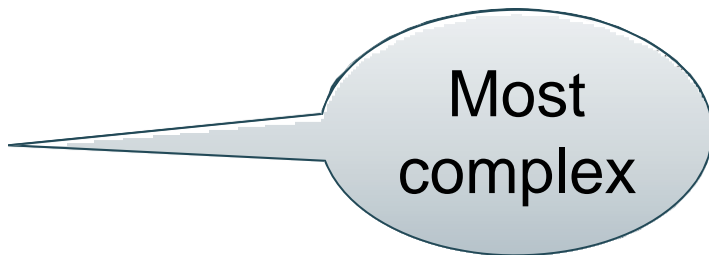
# Constraints in SQL

Constraints in SQL:

- Keys, foreign keys
- Attribute-level constraints
- Tuple-level constraints
- Global constraints: assertions



simplest



Most  
complex

- The more complex the constraint, the harder it is to check and to enforce...
  - (Still, performance is secondary to correctness.)



# Key Constraints

Product(name, category)

```
CREATE TABLE Product (  
    name CHAR(30) PRIMARY KEY,  
    category VARCHAR(20))
```

OR:

```
CREATE TABLE Product (  
    name CHAR(30),  
    category  
    VARCHAR(20),  
    PRIMARY KEY (name))
```

# Keys with Multiple Attributes

Product(name, category,  
price)

```
CREATE TABLE Product (  
    name CHAR(30),  
    category VARCHAR(20),  
    price INT,  
    PRIMARY KEY (name, category))
```

Name	Category	Price
Gizmo	Gadget	10
Camera	Photo	20
Gizmo	Photo	30
<del>Gizmo</del>	<del>Gadget</del>	<del>40</del>

# Other Keys

```
CREATE TABLE Product (  
    productID    CHAR(10),  
    name CHAR(30),  
    category VARCHAR(20),  
    price INT,  
    PRIMARY KEY (productID),  
    UNIQUE (name, category))
```

There is at most one **PRIMARY KEY**; there can be many **UNIQUE**

# Foreign Key Constraints

```
CREATE TABLE Purchase (  
    prodName CHAR(30)  
    REFERENCES Product(name),  
    date DATETIME)
```

Referential  
integrity  
constraints

prodName is a **foreign key** to Product(name)

name must be a **key** in Product  
(i.e., PK or UNIQUE)

May write  
just Product  
if name is PK

# Foreign Key Constraints

- Example with multi-attribute primary key

```
CREATE TABLE Purchase (  
    prodName CHAR(30),  
    category  
    VARCHAR(20), date  
    DATETIME,  
    FOREIGN KEY (prodName, category)  
    REFERENCES Product(name, category))
```

- (name, category) must be a KEY in Product

# What happens when data changes?

Types of updates:

- In Purchase: insert/update
- In Product: delete/update

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Store
Gizmo	Wiz
Camera	Ritz
Camera	Wiz

# What happens when data changes?

SQL has three options for maintaining referential integrity on changes:

- RESTRICT not allowed
- NO ACTION reject bad modifications (default)
- CASCADE after delete/update do delete/update
- SET NULL set foreign-key field to NULL
- SET DEFAULT set FK field to default value
  - need to be declared with column, e.g.,  
CREATE TABLE Product (pid INT DEFAULT 42)

# Maintaining Referential Integrity

```
CREATE TABLE Purchase (  
    prodName CHAR(30),  
    category  
    VARCHAR(20), date  
    DATETIME,  
    FOREIGN KEY (prodName, category)  
    REFERENCES Product(name,  
    category) ON UPDATE CASCADE  
    ON DELETE SET NULL)
```

Product

Name	Category
Gizmo	gadget
Camera	Photo
OneClick	Photo

Purchase

ProdName	Category
Gizmo	gadget
Snap	Camera
EasyShoot	Camera



# Constraints on Attributes and Tuples

- Constraints on attributes:
  - NOT NULL** -- obvious meaning...
  - CHECK** condition -- any condition !
- Constraints on tuples
  - CHECK** condition

# Constraints on Attributes and Tuples

```
CREATE TABLE Product (  
    productID    CHAR(10),  
    name CHAR(30),  
    category VARCHAR(20),  
    price INT CHECK (price > 0),  
    PRIMARY KEY (productID))
```

# Constraints on Attributes and Tuples

```
CREATE TABLE Product (  
    productID    CHAR(10),  
    name CHAR(30),  
    category VARCHAR(20)  
        CHECK (category in ('toy', 'gadget', 'apparel')),  
    price INT CHECK (price > 0),  
    PRIMARY KEY (productID))
```

# Constraints on Attributes and Tuples

```
CREATE TABLE Product (  
    productID  
        CHAR(10),  
    name CHAR(30) NOT NULL,  
    category VARCHAR(20)  
        CHECK (category in ('toy', 'gadget', 'apparel')),  
    price INT CHECK (price > 0),  
    PRIMARY KEY (productID))
```

# Constraints on Attributes and Tuples

```
CREATE TABLE R (  
    A int NOT NULL,  
    B int CHECK (B > 50 and B < 100),  
    C varchar(20),  
    D int,  
    CHECK (C >= 'd' or D > 0))
```

# Constraints on Attributes and Tuples

What does this constraint do?

```
CREATE TABLE Purchase
( prodName
  CHAR(30)
  CHECK (prodName IN
        (SELECT Product.name FROM
         Product)),
  date DATETIME NOT NULL)
```

What is the difference from Foreign Key?

# General Assertions

```
CREATE ASSERTION myAssert  
  CHECK (NOT EXISTS(  
    SELECT Product.name  
    FROM Product, Purchase  
    WHERE Product.name = Purchase.prodName  
    GROUP BY Product.name  
    HAVING count(*) > 200) )
```

But most DBMSs do not implement assertions,  
because it is hard to support them efficiently.  
Instead, DBMSs provide triggers