

Introduction to Database Systems

NoSQL & JSON
(mostly not in textbook...
only Ch 11.1)


NoSQL Motivation

- Originally motivated by Web 2.0 applications
- Goal is to scale simple OLTP-style workloads to **millions or billions of users**
 - Ex: Facebook has 1.2B *daily* active users
 - use often correlated in time in each region
 - > 10M req/sec if 25% of users arrive within one hour
 - SQL Server would collapse under that workload
- Users are doing both updates and reads

What is the Problem?

- Single server DBMS are too small for Web data
- **Solution:** scale out to multiple servers
- This is hard for the *entire* functionality of DBMS
 - as we will see next...
- NoSQL: reduce functionality for easier scaling
 - Simpler data model
 - Fewer guarantees

Serverless Architecture




Client-Server Architecture

Supports many apps and many users simultaneously

Server Machine

Client Applications



Connection (JDBC, ODBC)




- One server running the database
- Many clients, connecting via the ODBC or JDBC (Java Database Connectivity) protocol


Client-Server

- One *server* that runs the DBMS (or RDBMS):
 - Your own desktop, or
 - Some beefy system, or
 - A cloud service (SQL Azure)
- Many *clients* run apps and connect to DBMS
 - Microsoft's Management Studio (for SQL Server), or
 - psql (for postgres)
 - Some Java program (HW8) or some C++ program
- Clients “talk” to server using JDBC/ODBC protocol

3-Tiered Architecture




3-Tiered Architecture






3-Tiered Architecture

Replicate
App server for
scale up



Connection
(e.g., JDBC)



HTTP/SSL

Why don't we replicate
the DB server too?




Replicating the Database

- Much harder because the state must be unique. In other words, the database must act as a whole
 - Current DB instance must always be *consistent*
 - Ex: foreign keys must exist
 - as a result, some updates must occur simultaneously
- Two basic approaches:
 - Scale up through [partitioning](#)
 - Scale up through [replication](#)


Scale Through Partitioning

- Partition the database across many machines in a cluster
 - Database could fit in main memory
 - Queries spread across these machines
- Can increase throughput
- Easy for (simple) writes but reads become harder



Scale Through Replication

- Create multiple copies of each database partition
- Spread queries across these replicas
- Can increase throughput and lower latency
- Can also improve fault-tolerance
- Easy for reads but writes become harder



NoSQL Data Models

Taxonomy based on data models:

- ☞ • **Key-value stores**
 - e.g., Project Voldemort, Memcached
- **Document stores**
 - e.g., SimpleDB, CouchDB, MongoDB
- **Extensible Record Stores**
 - e.g., HBase, Cassandra, PNUTS

Key-Value Stores Features

- **Data model:** (key, value) pairs
 - Key = string/integer, unique for the entire data
 - Value = can be anything (very complex object)
- **Operations**
 - Get(key), Put(key, value)
 - Operations on value not supported
- **Distribution / Partitioning**
 - No replication: key k is stored at server $h(k)$
 - 3-way replication: key k is stored at $h1(k)$, $h2(k)$, $h3(k)$

How does get(k) work? How does put(k , v) work?

Flights(fid, date, carrier, flight_num, origin, dest, ...)
Carriers(cid, name)

Example

- How would you represent the Flights data as (key, value) pairs?
- Option 1: key=fid, value=entire flight record
- Option 2: key=date, value=all flights that day
- Option 3: key=(origin, dest), value=all flights between


How does query processing work?

Key-Value Stores Internals

- Data remains in main memory
 - one implementation: distributed hash table
- Most systems also offer a persistence option
- Others use replication to provide fault-tolerance
 - Asynchronous or synchronous replication
 - Tunable consistency: read/write one replica or majority
- Some offer transactions, others do not
 - multi-version concurrency control or locking
- No secondary indices

Data Models

Taxonomy based on data models:

- Key-value stores
 - e.g., Project Voldemort, Memcached
-  • Document stores
 - e.g., SimpleDB, CouchDB, MongoDB
- Extensible Record Stores
 - e.g., HBase, Cassandra, PNUTS


Document Stores Features

- **Data model:** (key, document) pairs
 - Key = string/integer, unique for the entire data
 - Document = JSON or XML
- **Operations**
 - Get/put document by key
 - Limited, non-standard query language on JSON
- **Distribution / Partitioning**
 - Entire documents, as for key/value pairs

We will discuss JSon today or tomorrow

Data Models

Taxonomy based on data models:

- Key-value stores
 - e.g., Project Voldemort, Memcached
- Document stores
 - e.g., SimpleDB, CouchDB, MongoDB
-  • Extensible Record Stores
 - e.g., HBase, Cassandra, PNUTS

Extensible Record Stores

- Based on Google's BigTable
 - HBase is an open source implementation of BigTable
- Data model is rows and columns
 - can add both new rows and new columns
- Scalability by splitting rows and columns over nodes
 - Rows partitioned through hashing on primary key
 - Columns of a table are distributed over multiple nodes by using “column groups”

NoSQL Summary

- Simpler data model with weaker guarantees
- But they scale as far as we need them to
- Meanwhile...
SQL systems continue to improve

Recent SQL Progress

- Modern systems need to store data across the globe
 - individual data centers go offline
 - need servers close to users to be efficient
- Speed of light is a fundamental limit
 - 200+ms latency (across US) is visible to users
- Systems must weaken guarantees
- Google's Spanner (supports SQL):
 - write data over the whole globe (a bit slowly)
 - reads occur slightly in the *past*

Prediction

- My guess: SQL will win again
- Pieces are out there already
 - Spanner: multi-node transactions
 - AsterixDB: multi-node query optimization
- For now, NoSQL still offers key benefits

JSon and Semi-structured Data

Where We Are

- So far we have studied the relational data model
 - Data is stored in tables (relations)
 - Queries are expressions in the SQL / Datalog / relational algebra
- Today: Semi-structured data model
 - Popular formats today: XML, JSon, protobuf

JSON

- 10 years ago...
 - JavaScript interpreters were very slow
 - native browser function parsed JSON 100x faster
- XML was also an option, but
 - IE had a memory leak in its XML parser
- JSON used in Gmail etc. for this reason
- Spread organically to server-side systems

JSON - Overview

- JavaScript Object Notation = lightweight text-based open standard designed for human-readable data interchange. Interfaces in C, C++, Java, Python, Perl, etc.
- The filename extension is .json.

We will emphasize JSon as semi-structured data

JSON vs. Relational

- Relational data model
 - Rigid flat structure (tables)
 - Schema must be fixed in advanced
 - Binary representation: good for performance, bad for exchange
 - Query language based on Relational Calculus
- Semi-structured data model / JSON
 - Flexible, nested structure (trees)
 - Does not require predefined schema ("self describing")
 - Text representation: good for exchange, bad for performance
 - Most common use: Language API; query languages emerging

Json Syntax

```
{ "book": [  
  {"id": "01",  
    "language": "Java",  
    "author": "H. Javeson",  
    "year": 2015  
  },  
  {"id": "07",  
    "language": "C++",  
    "edition": "second"  
    "author": "E. Sepp",  
    "price": 22.25  
  }  
]  
}
```

JSON Terminology

- Curly braces hold objects
 - Each object is a list of name/value pairs separated by , (comma)
 - Each pair is a name is followed by ':' (colon) followed by the value
- Square brackets hold arrays and values are separated by , (comma).
- Data made up of objects, lists, and atomic values (integers, floats, strings, booleans).


Json Data Structures

- Collections of name-value pairs:
 - {“name1”: value1, “name2”: value2, ...}
 - The “name” is also called a “key”
- Ordered lists of values:
 - [obj1, obj2, obj3, ...]

Avoid Using Duplicate Keys

The standard allows them, but many implementations don't

```
{  
  "id": "07",  
  "title": "Databases",  
  "author": "Garcia-Molina",  
  "author": "Ullman",  
  "author": "Widom"  
}
```



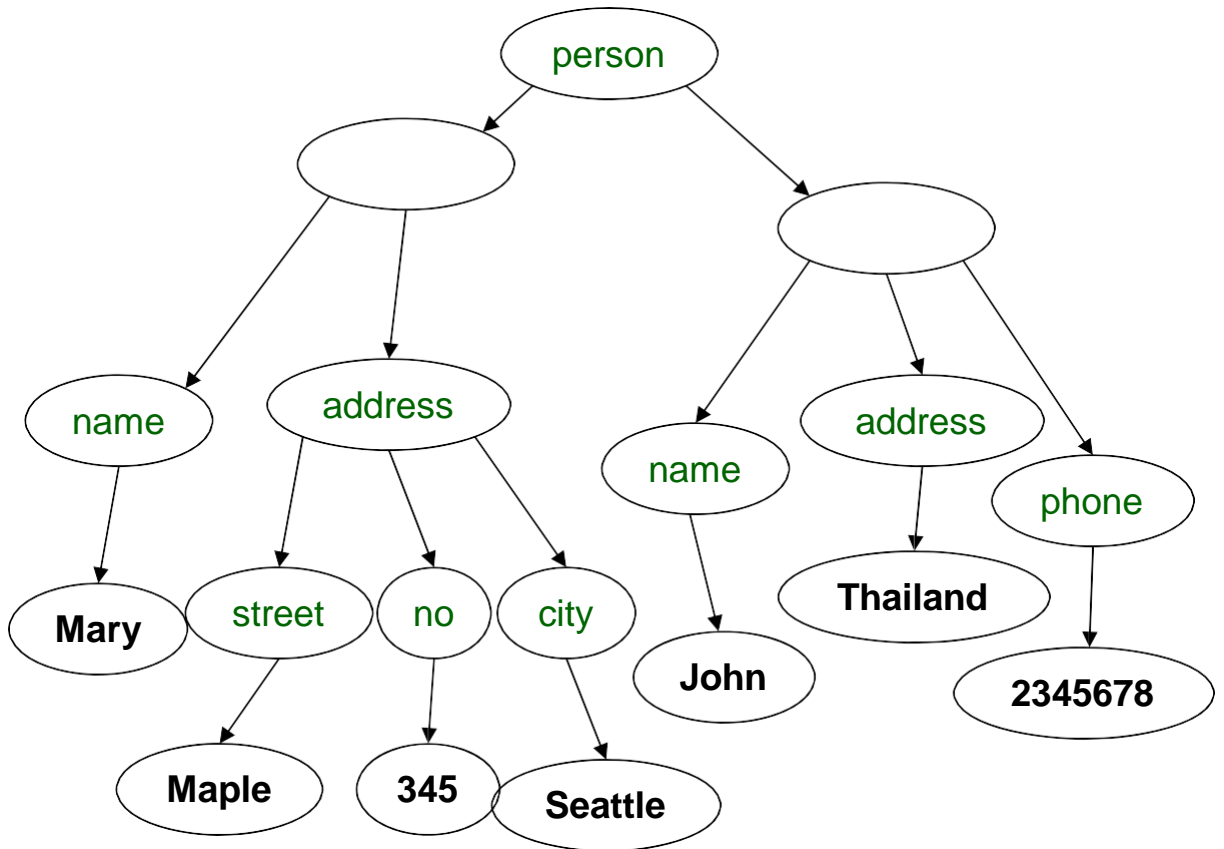
```
{  
  "id": "07",  
  "title": "Databases",  
  "author": ["Garcia-Molina",  
             "Ullman",  
             "Widom"]  
}
```


JSon Data Types

- Number
- String = double-quoted
- Boolean = true or false
- null / empty

JSON Semantics: a Tree !

```
{  
  "person":  
    [  
      {  
        "name": "Mary",  
        "address":  
          {  
            "street": "Maple",  
            "no": 345,  
            "city": "Seattle"  
          }  
      },  
      {  
        "name": "John",  
        "address": "Thailand",  
        "phone": 2345678  
      }  
    ]  
}
```




JSon Data

- JSon is **self-describing**
- Schema elements become part of the data
 - Relational schema: **person(name, phone)**
 - In JSon “**person**”, “**name**”, “**phone**” are part of the data, and are repeated many times
- Consequence: JSon is much more flexible
 - also uses more space (but can be compressed)
- JSon is an example of **semi-structured** data

Mapping Relational Data to JSON

Person

name	phone
John	3634
Sue	6343
Dirk	6363



```
{  
  "person": [  
    {  
      "name": "John",  
      "phone": 3634  
    },  
    {  
      "name": "Sue",  
      "phone": 6343  
    },  
    {  
      "name": "Dirk",  
      "phone": 6383  
    }  
  ]  
}
```

Mapping Relational Data to JSON

May inline foreign keys

Person

name	phone
John	3634
Sue	6343

Orders


personName	date	product
John	2002	Gizmo
John	2004	Gadget
Sue	2002	Gadget

```
{
  "Person": [
    {
      "name": "John",
      "phone": 3646,
      "Orders": [
        {
          "date": 2002,
          "product": "Gizmo"
        },
        {
          "date": 2004,
          "product": "Gadget"
        }
      ]
    },
    {
      "name": "Sue",
      "phone": 6343,
      "Orders": [
        {
          "date": 2002,
          "product": "Gadget"
        }
      ]
    }
  ]
}
```

JSON = Semi-structured Data (1/3)

- Missing attributes:

```
{  
  "person":  
    [  
      {"name": "John", "phone": 1234},  
      {"name": "Joe"}  
    ]  
}
```



- Could represent a table with nulls

name	phone
John	1234
Joe	-

JSON = Semi-structured Data (2/3)

- Repeated attributes

```
{  
  "person":  
    [{  
      "name": "John", "phone": 1234,  
      "name": "Mary", "phone": [1234, 5678]}]  
}
```

Two phones !

- Impossible in one table:

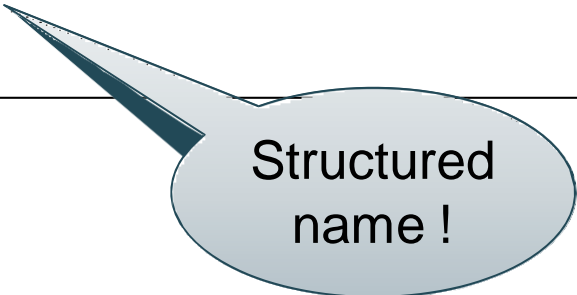
name	phone	
Mary	2345	3456

???

JSON = Semi-structured Data (3/3)

- Attributes with different types in different objects

```
{  
  "person": [  
    {  
      "name": "Sue",  
      "phone": 3456  
    },  
    {  
      "name": {  
        "first": "John",  
        "last": "Smith"  
      },  
      "phone": 2345  
    }  
  ]  
}
```



Structured
name !

- Nested collections
- Heterogeneous collections

Discussion

- *Data exchange formats*
 - well suited for exchanging data between apps
 - XML, JSON, Protobuf
- Increasingly, some systems use them as a data model:
 - SQL Server: supports for XML-valued relations
 - CouchBase, MongoDB: JSON as data model
 - Dremel (BigQuery): Protobuf as data model

Query Languages for Semi-Structured Data

- XML: XPath, XQuery (see end of lecture, textbook)
 - Supported inside many RDBMS (SQL Server, DB2, Oracle)
 - Several standalone XPath/XQuery engines
- Protobuf: used internally by Google, and externally in BigQuery. similar to compiled JSON
- JSON:
 - CouchBase: N1QL
 - MongoDB: has a pattern-based language
 - JSONiq <http://www.jsoniq.org/>
 - AsterixDB: AQL and SQL++